

Skin Disease Detection Using Deep Learning

Submitted in partial fulfillment of the requirements for the award of the degree of

BACHELOR OF TECHNOLOGY

by

Akshit Rawat Sec B(2018130)

Deepak Singh Sec B(2018304)

Rohit Kumar Sec A(2018664)

Dipin Bhandari Sec B(2018324)

To

Dr. Satvik Vats

Associate Professor

Department Of Computer Science & Engineering



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
GRAPHIC ERA HILL UNIVERSITY, DEHRADUN
DEHRADUN – 248002 (INDIA)**

CERTIFICATE

We **Akshit Rawat** , **Deepak Singh**, **Rohit Kumar** and **Dipin Bhandari**, students of **B.Tech CSE VIII Semester**, Department of Computer Science and Engineering, Graphic Era Hill University, Dehradun, declare that the technical project work entitled “Comparison of Performance of Pretrained Model and Customized CNN on Skin Disease Using Deep Learning” has been carried out by us and submitting in partial fulfillment of the course requirements for the award of degree in Bachelor Of Technology of **Graphic Era Hill University**, Dehradun during the academic year **2023-2024**. This synopsis has not been submitted to any other university for the award of any other degree or diploma.

Place: Dehradun

Date: 22 May 2024

ACKNOWLEDGEMENT

Hereby we are submitting the project report on “Comparison of Performance of Pretrained Model and Customized CNN on Skin Disease Using Deep Learning”, as per the scheme of Graphic Era Hill University, Dehradun. We express the deepest sense of gratitude to Mr. Satvik Vats, Asst. Professor, Department of Computer Science and Engineering for the invaluable guidance extended at every stage and in every possible way. I am very much thankful to all the faculty members of the Department of Computer Science and Engineering, friends and my parents for their constant encouragement, support and help throughout the period of project conduction.

Akshit Rawat

Univ. Roll No : 2018130

Sec/Class roll no. : B /8

Deepak Singh

Univ. Roll No : 2018304

Sec/ Class roll no.: B/24

Rohit Kumar

Univ. Roll No : 2018664

Sec/Class roll no. : A/42

Dipin Bhandri

Univ. Roll No : 2018324

Sec/Class roll no. : B/25

ABSTRACT

Accurately identifying skin conditions is essential for providing appropriate medical therapy. The intricate and often subtle patterns presented by various skin disorders can pose significant challenges for diagnosis. These complexities necessitate the development of advanced and precise diagnostic tools. Our study embarks on an in-depth exploration of utilizing both pretrained models and custom models to enhance the prediction of skin disorders, leveraging the strengths of each approach to achieve superior diagnostic accuracy.

Pretrained models, such as Inception-V3 and VGG-16, have demonstrated remarkable capabilities in general image recognition tasks. These models are equipped with deep learning architectures that have been pre-trained on extensive datasets, allowing them to extract and learn from a wide array of image features. By applying transfer learning, we can repurpose these models for specific tasks in dermatology..

While pretrained models offer substantial benefits, custom CNNs provide a tailored approach specifically designed to address the unique characteristics of skin images. Custom CNN layers are meticulously crafted to focus on the specific features pertinent to dermatological analysis, such as color variations, texture patterns, and lesion shapes. By constructing a custom architecture, we ensure that the model is optimally aligned with the nuances of skin disorder detection.

The HAM10000 dataset serves as the cornerstone of our study, providing a comprehensive and varied set of dermoscopic images necessary for training and evaluation. This dataset includes images of various skin conditions, ranging from benign lesions to malignant melanomas, thereby offering a robust platform for model training.

To ensure rigorous evaluation, we employ a variety of metrics, including accuracy, precision, recall, and F1-score. These metrics provide a comprehensive assessment of the models' performance, enabling us to compare their effectiveness in predicting skin disorders.

In summary, Our study represents a significant advancement in the field of automated dermatological diagnosis and medical image analysis. By leveraging the strengths of both pretrained models and custom CNNs, we strive to develop a robust and reliable diagnostic tool that can assist clinicians in accurately identifying skin conditions.

INDEX

ACKNOWLEDGMENT

ABSTRACT

LIST OF TABLES

LIST OF FIGURES

ABBREVIATIONS

NOTATIONS

1. INTRODUCTION

- 1.1 Historical Perspective on Skin Disease Diagnosis**
- 1.2 The Complexity of Skin Disease Prediction**
- 1.3 Leveraging Advanced AI Techniques**
- 1.4 Artificial Neural Networks**
- 1.5 Project Objectives and Scope**
- 1.6 Methodology and Approach**
- 1.7 Expected Outcomes and Impact**

2. LITERATURE SURVEY

3. METHODOLOGY

- 3.1 Dataset Description**
- 3.2 VGG16 Model**
- 3.3 InceptionV3 Model**
- 3.4 Custom CNN Model**
- 3.5 Evaluation**

4. RESULTS AND ANALYSIS

- 4.1 Quantitative Analysis**
- 4.2 Qualitative Analysis**
- 4.3 comparison Analysis**

5. CONCLUSION AND FUTURE SCOPE

APPENDIX

REFERENCES

LIST OF FIGURES

Figure 1: Flow Chart

Figure 3.1 : Pie Chart of Imbalance Data of Ham 10000

Figure 3.2 : Images of Ham10000 Dataset

Figure 3.3: Architecture of VGG16

Figure 3.4 Visualizing K-Fold Data Splitting

Figure 3.5 Training and loss Graph of Vgg-16

Figure 3.6: The Architecture of Inception v3

Figure 3.7 : Stem Block of Inception V3

Figure 3.8 : Inception A-Block

Figure 3.9 Inception B Block

Figure 3.10 Accuracy Graph of the Inception V3

Figure 3.11 Loss Graph of the Inception V3

Figure 3.12 Architecture of Custom CNN Model

Figure 3.13 Accuracy of the Custom Model

Figure 3.14 Loss Graph of Custom Model

ABBREVIATIONS

- CNN: Convolutional Neural Network
- ReLU: Rectified Linear Unit
- RGB: Red, Green, Blue (color model)
- VGG: Visual Geometry Group
- IEEE: Institute of Electrical and Electronics Engineers
- HAM10000 : Human Against Machine with 10000 training images
- MEL : Melanoma
- BKL : Benign Keratosis
- NV : Melanocytic Nevus
- Adam : Adaptive Moment Estimation
- Conv : Convolutional Layer

CHAPTER 1

INTRODUCTION

The accurate and timely prediction of skin diseases has long been a critical goal in the field of dermatology and medical diagnostics. Traditional diagnostic methods relied heavily on the expertise of dermatologists, requiring them to meticulously examine and diagnose skin conditions—a process that is not only labor-intensive but also subject to human error and variability. The advent of deep learning and artificial intelligence has transformed this challenging domain, enabling the development of sophisticated models that can automatically analyze and predict skin diseases with remarkable accuracy and consistency. This project delves into the evolving landscape of skin disease prediction, providing an in-depth exploration of the transition from manual diagnosis to cutting-edge AI-driven solutions.

1.1 Historical Perspective on Skin Disease Diagnosis

Skin diseases, ranging from common conditions like eczema and psoriasis to serious ailments such as melanoma, have long been documented throughout medical history. Early methods of diagnosis were largely based on visual inspections and the expertise of trained physicians. These diagnostic images and records serve as crucial medical documentation, providing insights into the prevalence, treatment, and progression of various skin conditions over time. However, the manual nature of these diagnoses often led to inconsistencies and diagnostic delays.

Throughout history, skin conditions have been described and categorized based on their visual manifestations. Ancient texts and medical manuscripts provide detailed accounts of various skin ailments and their treatments. With the advent of photography and more advanced imaging techniques, the ability to document and study these conditions improved significantly. Despite these advancements, the primary tool for diagnosis remained the trained human eye, which, while effective, was not infallible. The variability in human judgment and the subjective nature of

visual assessments often resulted in diagnostic errors and missed early detections, particularly in cases of rare or atypical presentations of common diseases.

1.2 The Complexity of Skin Disease Prediction

The prediction and diagnosis of skin diseases pose significant challenges due to the vast diversity of skin conditions, their overlapping symptoms, and the subtle differences that can distinguish one condition from another. Traditional diagnostic approaches, while effective, are hindered by their reliance on subjective visual assessments and the limited availability of dermatological expertise. This underscores the necessity for advanced, automated solutions that can enhance diagnostic accuracy, efficiency, and accessibility.

Skin diseases encompass a wide spectrum of conditions, each with its own set of characteristics and clinical presentations. Some conditions, such as acne and dermatitis, are widespread and well-studied, while others, like certain types of skin cancer, are less common but far more dangerous. The clinical manifestations of these diseases can vary widely, even among patients with the same condition, making accurate diagnosis particularly challenging. Additionally, environmental factors, genetic predispositions, and individual patient histories contribute to the complexity of diagnosis. Dermatologists must consider all these variables, often under the constraints of limited consultation time and high patient volumes, which can compromise the thoroughness and accuracy of their assessments.

1.3 Leveraging Advanced AI Techniques

Recent advancements in artificial intelligence, particularly in the field of deep learning, have revolutionized medical imaging and diagnostic processes. Convolutional Neural Networks (CNNs), a powerful type of deep learning architecture, have demonstrated exceptional performance in image recognition tasks, including the prediction of skin diseases. By learning from extensive datasets of annotated skin images, CNNs can accurately identify and classify a wide range of skin conditions. This capability opens up new possibilities for automated, reliable, and scalable skin disease prediction, significantly augmenting traditional diagnostic practices. The power of CNNs lies in their ability to learn hierarchical

features from input images. These networks consist of multiple layers that progressively extract more abstract and complex features from the raw pixel data. In the context of skin disease prediction, CNNs can be trained on large datasets comprising thousands of images of various skin conditions, each labeled with its corresponding diagnosis. Through this training process, CNNs learn to recognize intricate patterns and features that may be indicative of specific diseases, even those that are subtle and not easily discernible by the human eye. The deployment of these models in clinical settings can assist dermatologists by providing a second opinion, prioritizing cases that require urgent attention, and expanding diagnostic capabilities to regions with limited access to dermatological expertise. As a result, the integration of AI-driven solutions in dermatology promises to enhance patient outcomes, streamline diagnostic workflows, and contribute to the broader field of medical research and education.

1.4 Artificial neural networks

Artificial neural networks are built on the principles of the structure and operation of human neurons. It is also known as neural networks or neural nets. An artificial neural network's input layer, which is the first layer, receives input from external sources and passes it on to the hidden layer, which is the second layer. Each neuron in the hidden layer gets information from the neurons in the previous layer, computes the weighted total, and then transfers it to the neurons in the next layer. These connections are weighted, which means that the impacts of the inputs from the preceding layer are more or less optimized by giving each input a distinct weight. These weights are then adjusted during the training process to enhance the performance of the model.

Artificial neurons, also known as units, are found in artificial neural networks. The whole Artificial Neural Network is composed of these artificial neurons, which are arranged in a series of layers. The complexities of neural networks will depend on the complexities of the underlying patterns in the dataset whether a layer has a dozen units or millions of units. Commonly, Artificial Neural Networks have an input layer, an output layer as well as hidden layers. The input layer receives data from the outside world which the neural network needs to analyze or learn about.

In a fully connected artificial neural network, there is an input layer and one or more hidden layers connected one after the other. Each neuron receives input from the previous layer neurons or the input layer. The output of one neuron becomes the input to other neurons in the next layer of the network, and this process continues until the final layer produces the output of the network. Then, after passing through one or more hidden layers, this data is transformed into valuable data for the output layer. Finally, the output layer provides an output in the form of an artificial neural network's response to the data that comes in.

Units are linked to one another from one layer to another in the bulk of neural networks. Each of these links has weights that control how much one unit influences another. The neural network learns more and more about the data as it moves from one unit to another, ultimately producing an output from the output layer.

1.5 Project Objectives and Scope

The overarching objective of our project is to leverage CNN-based techniques to develop a sophisticated skin disease prediction application tailored specifically for dermatological diagnostics. This project aims to achieve several critical goals. First, we intend to develop a robust CNN model trained on curated datasets of annotated skin images representing a wide variety of skin conditions. This model will be designed to accurately identify and classify skin diseases by learning from extensive, diverse datasets, ensuring that it can handle a broad spectrum of conditions and patient demographics.

Second, we will implement an intuitive and user-friendly interface that facilitates interaction with the skin disease prediction application. This interface will be designed to streamline the diagnostic process, allowing users to easily upload images, receive predictions, and access detailed information about the diagnosed conditions. The emphasis on usability aims to make the tool accessible to both healthcare professionals and patients, enhancing its practical utility in real-world settings.

Third, we will incorporate mechanisms for users to provide feedback and contribute to the improvement of prediction algorithms. User feedback will be vital for refining the model, addressing any inaccuracies, and ensuring that the predictions remain reliable and up-to-date. This feedback loop will help the system learn and evolve continuously, adapting to new data and emerging skin conditions.

Fourth, we will explore and compare the performance of pretrained models, including VGG-16 and Inception V3, against our customized CNN model. By benchmarking these models, we aim to understand the strengths and weaknesses of each approach, providing insights into which architectures are most effective for skin disease prediction. This comparative analysis will be crucial for optimizing our model and ensuring it meets the highest standards of accuracy and efficiency.

Finally, we will evaluate the accuracy, efficiency, and reliability of our models by benchmarking them against existing models in the field. This evaluation will involve rigorous testing and validation processes to ensure that our models perform well in various scenarios and on different types of data. By doing so, we can confidently deploy our model in clinical settings, knowing that it meets the necessary performance criteria.

1.6 Methodology and Approach

Our approach involves a comprehensive methodology that encompasses data collection, model development, user interface design, and evaluation metrics. We will collaborate with dermatologists, medical researchers, and domain experts to curate and preprocess the HAM10000 dataset, ensuring the accuracy and relevance of the skin disorder images. Model development will involve training and fine-tuning both pretrained models, such as VGG-16 and Inception-V3, and our custom-designed CNN architecture. By utilizing state-of-the-art techniques in deep learning and medical image analysis, we aim to create robust models capable of accurately predicting various skin conditions.

1.7 Expected Outcomes and Impact

At the conclusion of the project, we anticipate delivering a comprehensive comparative analysis of the performance and accuracy of pretrained models versus our custom CNN model on the HAM10000 dataset. This analysis will serve as a valuable resource for the medical community, providing insights into the most effective approaches for automated skin disease prediction. The resulting models and findings will not only assist clinicians in improving diagnostic accuracy but also contribute to the broader field of medical AI research. Furthermore, our project aims to advance discussions on the ethical and societal implications of AI-driven technologies in healthcare, emphasizing the importance of transparency, interpretability, and responsible usage in medical diagnostics.

CHAPTER 2

LITERATURE SURVEY

Viswanatha Reddy Allu Gunti (2021) This study addresses melanoma, a highly dangerous skin cancer, particularly prevalent among white populations in Australia and New Zealand. The research proposes a deep learning technique using a Convolutional Neural Network (CNN) to differentiate between melanoma types—malignant, superficial spreading, and nodular. Utilizing data from dermnetnz.org, the CNN classifier surpasses existing methods in diagnostic accuracy, emphasizing the importance of early detection for effective treatment.

Nawal Soliman AL Kolifi AL Enezi (2019) This research focuses on the significant health concern posed by skin diseases in Saudi Arabia due to its harsh climate. It introduces a cost-effective image processing method for rapid and accurate detection of skin diseases. The method uses digital images and a pretrained CNN for feature extraction, followed by Multiclass SVM for classification. The system achieves 100% accuracy in detecting three specific skin diseases, highlighting the potential of image processing techniques in dermatological screening in high-incidence regions.

P. Nagaraj, V. Muneeswaran, K. Jaya Krishna, K. Yerriswamy Reddy, J. Rock Morries, G. Pavan Kumar (2019) This study explores the application of CNNs in diagnosing common but painful skin disorders such as chickenpox, impetigo, and scabies, which are prevalent and costly to treat. By processing images of affected skin areas through a CNN, the system provides precise diagnostic results, aiding in the timely identification and management of these conditions.

Rola EL SALEH, Sambit BAKHSHI, Amine NAIT-ALI (December 2019) The paper suggests a method for automatic facial skin disease detection using a pre-trained deep CNN. By enhancing and resizing images, the model is trained to identify eight different facial skin diseases and normal skin with an accuracy rate of 88%. This approach could significantly speed up and improve the accuracy of

dermatological diagnosis.

Nazia Hameed, Antesar M. Shabut, M. A. Hossain (February 2019) This research underscores the global prevalence of skin diseases and their impact. It introduces a hybrid diagnostic method combining a deep CNN and error-correcting output code SVM to classify skin lesions into five categories. Using the AlexNet model on 9,144 images, the method achieves an accuracy of 86.21%, demonstrating the efficacy of this combined approach in dermatological diagnostics.

Muhammad Naseer Bajwa (March 2020) This study leverages deep learning to diagnose skin disorders, training Deep Neural Networks on large datasets from DermNet and the ISIC Archive. Using disease taxonomy for better classification, the model achieves 80% accuracy on DermNet and 93% on the ISIC Archive. The study highlights the effectiveness of deep learning in distinguishing between multiple skin conditions with high accuracy and reliability.

Li and Shen (2020) This study explored the use of pretrained models like GoogLeNet (Inception V3) for skin disease classification. By fine-tuning these models with large medical image datasets, they achieved significant improvements in diagnostic accuracy. The research emphasized the benefits of transfer learning, where knowledge from pretrained models can be leveraged to enhance the performance of skin disease classification systems

Esteva et al. (2017) Esteva and colleagues trained a CNN on a dataset of 129,450 clinical images covering over 2,000 diseases. The model's performance was on par with dermatologists in identifying skin cancer, illustrating the potential of deep learning to support medical professionals in diagnostic processes. This study underscored the power of deep learning in achieving high diagnostic accuracy for complex medical conditions

Han et al. (2018) Han and colleagues used a large dataset of 220,000 clinical images to train a deep CNN, achieving high accuracy in detecting 12 skin conditions. The study highlighted the importance of large and diverse datasets for training robust deep learning models, which are crucial for reliable medical diagnostics

Tschandl et al. This study compared various CNN architectures, including VGG-16 and InceptionV3 on the HAM10000 dataset. The results showed that while pretrained models offered strong baseline performance, custom models fine-tuned for specific skin conditions provided better specificity and sensitivity. This research highlighted the importance of model customization for enhanced diagnostic accuracy.

Brinker et al. (2019) Brinker and colleagues conducted a comparative study of several CNN architectures and found that ensembles of multiple models provided the best diagnostic performance for skin cancer detection. The study emphasized the value of combining different models to leverage their individual strengths, leading to improved overall accuracy.

CHAPTER 3

METHODOLOGY

Our research utilizes a customized convolutional neural network (CNN) alongside pre-trained CNN models, VGG16 and InceptionV3, to improve the prediction of skin illnesses. The primary goal is to develop a robust predictive model capable of accurately and consistently diagnosing a variety of skin conditions. The implementation of our models involved several key Python packages, including numpy, pandas, matplotlib.pyplot, sklearn.model_selection, sklearn.linear_model, and sklearn.metrics.

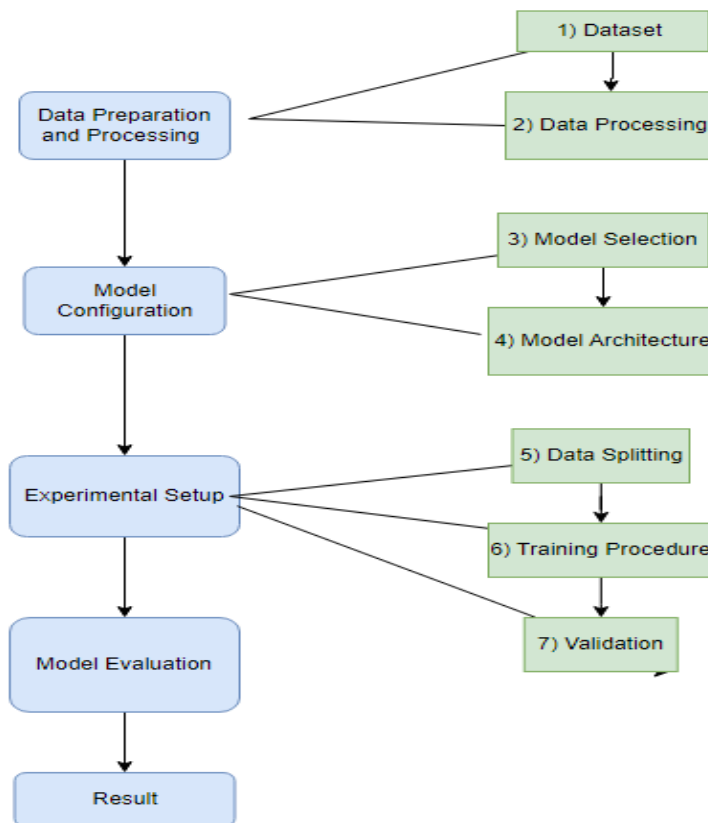


Figure 1: Flow Chart

Data Preparation and Processing :

We created a custom dataset inspired by the HAM10000 dataset, which contains 10,015 dermoscopic images. For our purposes, we classified the images into different classes by organizing them into seven subfolders, each named after a class label. From these, we selected three classes to form our custom dataset.

To increase the diversity within the dataset and enhance model robustness, we applied several image augmentation techniques using the image data generator function.

Rescaling involves normalizing pixel values to the range $[0, 1]$, which helps standardize the input data and facilitates the learning process. Shearing entails applying random shearing transformations, which distorts the image along a particular axis and aids in making the model more robust to variations in the data. Zooming includes randomly zooming into the images, allowing the model to focus on different parts of the image and enhancing its ability to recognize features at various scales. Brightness adjustment involves randomly altering the brightness levels of the images, simulating different lighting conditions and helping the model

Model Configuration :

We employed three different models for this study: InceptionV3 , VGG16, and a customized CNN model. VGG16 and InceptionV3 were selected due to their proven efficiency in handling deep neural networks with residual connections and their success in achieving high accuracy in image classification tasks through transfer learning.

1. Custom CNN Model :

Our customized CNN model features a hierarchical architecture with convolutional layers followed by max-pooling layers, enabling it to extract hierarchical features from input images. The Rectified Linear Unit (ReLU) activation function introduces non-linearity in hidden layers, aiding in capturing complex data relationships. The output layer facilitates multiclass classification, with three neurons representing skin condition classes. Softmax activation ensures accurate

classification by generating class probabilities. These design elements enhance our CNN model's effectiveness in diagnosing skin illnesses.

2. Fine-Tuning VGG16 and InceptionV3 Model :

In the initial stages, we utilized layer freezing to preserve pre-trained convolutional layers' features from models like VGG16 and InceptionV3. This allowed us to focus training efforts on new dense layers, enhancing efficiency. K-Fold cross-validation in VGG16 training improved model robustness, preventing overfitting. Additionally, we expanded our model by adding new dense layers with 256 neurons, using ReLU activation for feature extraction. The output layer, with three neurons, employed softmax activation for precise classification of skin conditions. We employed the Adam optimizer with varying learning rates to fine-tune model parameters, aiming to develop a highly accurate predictive model for diagnosing skin illnesses.

Experimental Setup :

To evaluate our models, we split our dataset into training, testing, and validation sets in a ratio of 7:2:1. This split ensures that we have sufficient data for training while also retaining enough data for unbiased testing and validation. The training set was used to train the models, providing a substantial amount of data for the models to learn from. The validation set was used to tune hyperparameters and prevent overfitting during training, allowing us to make necessary adjustments to improve model performance. The testing set was used to evaluate the final model performance, ensuring that the results reflect the model's ability to generalize to new, unseen data. Each model (VGG16, InceptionV3, and our customized CNN) was meticulously trained on the training set, validated on the validation set to fine-tune parameters, and finally tested on the testing set to assess their overall performance. This comprehensive evaluation approach helps ensure the robustness and reliability of our models in diagnosing skin conditions.

Model Evaluation :

We evaluated the performance of our models using several key metrics: training accuracy, validation accuracy, training loss, and validation loss. Training accuracy refers to the accuracy of the model on the training set, while validation accuracy measures the model's accuracy on the validation set. Similarly, training loss is the

loss incurred by the model on the training set, and validation loss is the loss incurred on the validation set. These metrics were crucial for monitoring the training process and ensuring that the models were learning effectively without overfitting. Additionally, we compared the performance of the different models to identify the best approach for diagnosing skin conditions. By structuring our methodology in this manner, we provide a comprehensive overview of the processes and techniques employed to develop and evaluate our predictive models for skin illness diagnosis. This structured approach ensures the reproducibility of our results and allows for the validation and improvement of our methods.

3.1 Dataset Description

The HAM10000 (Human Against Machine with 10000 training images) dataset is a comprehensive collection designed to support research in dermatology, particularly for the development and evaluation of automated skin lesion diagnosis systems. This dataset is meticulously curated to encompass a wide range of skin diseases, providing an invaluable resource for both academic research and practical applications in medical image analysis.

The HAM10000 dataset comprises 10,015 high-resolution labeled images, originally sized at 450x600 pixels. These images represent seven distinct classes of skin diseases: Melanoma (MEL), Melanocytic Nevus (NV), Basal Cell Carcinoma (BCC), Actinic Keratosis (AK), Benign Keratosis (BKL), Dermatofibroma (DF), and Vascular Lesion (VASC).

One of the notable challenges in the HAM10000 dataset is the significant class imbalance. More than 60% of the images belong to the "Melanocytic Nevus (NV)" class, while some classes are extremely rare, constituting less than 2% of the dataset. To address this imbalance and facilitate more robust model training and evaluation, the classes were reduced to three main categories, each comprising 1,099 images: Melanocytic Nevus (NV), Benign Keratosis (BKL), and Melanoma (MEL).

Prior to being used for model training, the images in the HAM10000 dataset underwent meticulous preprocessing to ensure consistency and quality. The preprocessing steps included resizing the images to 256x256 pixels for pretrained models to match the input size of models trained on ImageNet, and to 224x224 pixels for custom models. The images were normalized to a range of [0, 1] by

scaling pixel values to 1/255, and were processed in batches of 32 to optimize training efficiency.

The dataset was divided into three parts to facilitate comprehensive model evaluation: a training set comprising 2,506 images, a validation set of 657 images, and a testing set. This division allows for thorough training, fine-tuning, and evaluation of models, ensuring that they perform well not only on the training data but also on unseen images during validation and testing phases.

The HAM10000 dataset is pivotal for advancing research in automated skin lesion analysis. By providing a rich and diverse set of images, it enables researchers to develop and validate algorithms that can accurately identify and classify various skin diseases. The dataset's extensive coverage and detailed labeling allow for comprehensive training and testing, ensuring that models trained on this dataset are robust and effective in real-world medical applications.

Researchers leveraging the HAM10000 dataset can contribute to the development of diagnostic tools that enhance early detection and treatment of skin conditions, ultimately improving patient outcomes. Its detailed representation of multiple skin disease classes, despite the initial imbalance, offers a realistic foundation for tackling the challenges in dermatological image analysis, fostering innovations that can revolutionize skin disease diagnostics.

In essence, the HAM10000 dataset stands as an indispensable resource in the quest to advance medical imaging technologies, providing a robust platform for exploring the complexities and nuances of skin lesion classification.




Figure 3.1 : Pie Chart of Imbalance Data of Ham 10000




Figure 3.2 : Images of Ham10000 Dataset

3.2 VGG16 Model

Introduction to VGG16 :

VGG16 is a convolutional neural network (CNN) developed by the Visual Geometry Group (VGG) at the University of Oxford. Introduced by Karen Simonyan and Andrew Zisserman in their 2014 paper "Very Deep Convolutional Networks for Large-Scale Image Recognition," VGG16 was a major advancement in deep learning and image classification, significantly contributing to the development and understanding of CNN architectures.

Architecture of VGG16 :

VGG16 features an input layer for 224x224 pixel images with RGB color channels. It comprises 13 convolutional layers with 3x3 filters for detailed feature capture. Max-pooling layers of 2x2 size follow every two or three convolutional layers to reduce dimensions. The network then includes three fully connected layers: the first two with 4096 neurons each and the third with 1000 neurons for ImageNet classes. A softmax layer provides class probabilities, and dropout layers prevent overfitting. VGG16 boasts approximately 138 million parameters, making it notably large compared to earlier models.




Figure 3.3: Architecture of VGG16

Importance in Deep Learning :

1. Performance on ImageNet

VGG16 gained prominence by achieving top results in the ImageNet Large Scale Visual Recognition Challenge (ILSVRC) 2014. It achieved a top-5 error rate of 7.3%, showcasing its ability to classify images with high accuracy. This performance demonstrated the effectiveness of deeper networks and inspired further research into deeper architectures.

2. Depth and Simplicity

One of the key contributions of VGG16 was its demonstration that depth is a critical factor in the performance of CNNs. By using small (3x3) convolutional filters and stacking more layers, VGG16 showed that deep networks could learn more complex features. The simplicity of using uniform layer structures also made it easier to understand and implement.

3. Transfer Learning

In this approach, transfer learning plays a pivotal role, enabling us to harness the extensive training of VGG16 on the ImageNet dataset. By utilizing the pre-trained model as a feature extractor, we can capture fundamental visual patterns beneficial for a broad array of image classification tasks. Transfer learning offers several advantages: it notably reduces training time compared to starting from scratch, often leads to improved performance owing to the vast and diverse data the pre-trained models have been exposed to, and mitigates data requirements, enabling models to perform well even with smaller datasets.

Fine-tuning is another crucial step, involving the adjustment of weights in the pre-trained model to better suit the new dataset. Here, we selectively fine-tune the later layers of VGG16 while maintaining the earlier layers frozen. This selective fine-tuning approach involves freezing the early layers to preserve learned low-level features such as edges and textures, while training the later layers to adapt to the unique characteristics of the skin disease dataset, thus capturing nuanced patterns essential for effective classification. Adding custom layers atop the pre-trained VGG16 model allows us to tailor the network to our specific classification task.

These custom layers include batch normalization to enhance training stability and speed, dropout layers to prevent overfitting by promoting robust feature learning, dense layers with ReLU activation to introduce non-linearity and capture complex relationships within the data, and an output layer employing softmax activation for multi-class classification.

Model compilation involves selecting the Adam optimizer for its adaptive learning rate capabilities and efficiency, categorical cross entropy as the loss function suitable for multi-class classification problems, and accuracy as the primary evaluation metric to monitor the model's performance during training and validation. In practical application, the fine-tuned VGG16 model can be invaluable for skin disease classification, providing healthcare professionals with a potent tool to diagnose various skin conditions accurately.

By training the model on a labeled dataset of skin disease images, it can effectively distinguish between different conditions, offering reliable diagnostic support.

4. Foundation for Future Architectures

The VGG architecture influenced many subsequent models. For instance, the concept of deeper networks with small convolutional filters was adopted and extended in architectures like ResNet (Residual Networks) and Inception networks. The ideas from VGG16 also contributed to the development of efficient architectures designed for resource-constrained environments, such as MobileNet and SqueezeNet.

Application :

VGG16 has been widely used in various computer vision tasks beyond image classification:

Object Detection: By combining VGG16 with detection frameworks like R-CNN, Fast R-CNN, and Faster R-CNN, the architecture has been used to identify objects within images with high precision.

Semantic Segmentation: VGG16's feature maps have been utilized in segmentation models like Fully Convolutional Networks (FCNs) to assign a class label to each pixel in an image.

Image Retrieval: The deep features extracted by VGG16 have been used in image retrieval systems to find similar images within large databases.

Medical Imaging: In healthcare, VGG16 has been applied to tasks like tumor detection, skin lesion classification, and other diagnostic challenges, showcasing its versatility and robustness in various domains.

Conclusion :

VGG16's introduction marked a pivotal moment in the trajectory of deep learning and image classification. Its inception not only set new benchmarks but also ushered in an era of innovation, inspiring the development of more intricate and efficient models. By emphasizing the significance of depth in convolutional neural networks (CNNs), VGG16 played a crucial role in reshaping our understanding of network architectures and their capabilities in handling complex visual data. Moreover, its widespread adoption across diverse fields, from medical imaging to autonomous vehicles, underscores its enduring relevance and impact. As a result, VGG16 continues to serve as a cornerstone in both the theoretical exploration and practical application of deep learning principles in computer vision.

Fine-Tuned VGG16 Model for Custom Skin Disease Dataset :

1. Introduction

In the realm of deep learning, convolutional neural networks (CNNs) have become a cornerstone for image classification tasks. Among these, the VGG16 model, developed by the Visual Geometry Group at the University of Oxford, stands out due to its simplicity and effectiveness. This model has been widely adopted and fine-tuned for various specific tasks beyond its initial training on the ImageNet dataset. In this document, we delve into the detailed theory behind fine-tuning the VGG16 model for a custom skin disease dataset, examining each component of the code and its significance.

2. Loading the Pretrained VGG16 Model

The VGG16 model, a deep CNN trained on the ImageNet dataset containing over 14 million images across 1,000 categories, serves as a foundation for our task. Leveraging pretrained models allows us to tap into a wealth of learned feature representations, ranging from low-level details like edges to high-level semantics. By initializing the VGG16 model with pretrained weights, we expedite adaptation to new tasks with reduced data and computational requirements, a process known as transfer learning.

We excluded the top layers of the VGG16 network using the `include_top` parameter and set it to `false`. These top layers, designed for the original 1,000 ImageNet classes, are replaced with custom layers tailored to our specific task, such as skin disease classification. By retaining the convolutional base while excluding these top layers, we preserve the model's ability to capture generic visual features. The `input_shape` parameter specifies the dimensions of the input images in the custom dataset, ensuring compatibility with the network architecture. Specifically, width and height correspond to the image dimensions, while 3 represents the RGB color channels, ensuring correct processing of input images by the network and allowing our model to retain the convolutional base.

3. Setting Layers to Be Trainable

Initially, all layers of the VGG16 model are set to be trainable, allowing their weights to be updated during training. However, to strike a balance between leveraging pretrained knowledge and adapting to new data, a selective fine-tuning strategy is employed.

Our python code iterated through the layers of the VGG16 model, freezing the initial layers and only fine-tuning the later layers. This approach aims to preserve the valuable features learned from the ImageNet dataset while enabling the model to adapt to the specific features of the skin disease dataset. The layers before 'block5_conv1' are frozen, meaning their weights remain constant during training.

This strategy offers several benefits. Firstly, it preserves learned features by retaining the foundational knowledge acquired from ImageNet, enabling the model to recognize basic visual patterns effectively. Secondly, it helps in reducing overfitting by limiting the number of trainable parameters, which is particularly

advantageous when dealing with smaller datasets. Additionally, training fewer layers enhances computational efficiency, speeding up the fine-tuning process and reducing the computational load.

Selective training is integral to this approach, allowing only the later layers to be trainable. By doing so, the model can fine-tune high-level features specific to the new dataset while maintaining lower-level, generic features. This balance is critical for effective transfer learning, ensuring that the model adapts appropriately to the new task without sacrificing the valuable knowledge gained from pretraining on ImageNet.

4. Creating a Sequential Model

The Sequential model in Keras facilitates the creation of a linear stack of layers, providing a straightforward method for adding and configuring additional layers atop the base model. This modular approach simplifies both the construction and management of the model architecture.

By initializing a Sequential model and adding the VGG16 base model as the first layer, we establish a foundation upon which custom layers can be built to address our specific classification task. This integration enables the utilization of the pre-trained convolutional layers of VGG16 in conjunction with new layers designed to capture the distinctive features of the skin disease dataset.

5. Adding Custom Layers

We've enhanced the model by adding custom layers to further optimize its performance. First, we integrated Batch Normalization, a technique that stabilizes and accelerates the training process by normalizing the inputs of each mini-batch. This adjustment mitigates internal covariate shift and ensures that subsequent layer inputs are properly centered and scaled, facilitating quicker convergence and improved overall performance.

Next, we employed Dropout layers to prevent overfitting by randomly deactivating a portion of input units during training. By incorporating Dropout layers with a rate of 0.4, we ensured that 40% of neurons are randomly disabled during each training step. This strategy enhances the model's ability to generalize by reducing its reliance on specific network pathways.

Additionally, we included a Flatten layer to convert the 2D feature maps generated by the convolutional layers into a 1D feature vector. This transformation is essential before passing the data to fully connected (dense) layers, ensuring compatibility and efficient processing.

We then integrated a Dense layer with 256 neurons and ReLU activation to introduce non-linearity and enable the model to discern intricate patterns within the data. The weights of this dense layer were initialized using the He normal initializer for optimal training performance.

To further prevent overfitting, we appended an additional Dropout layer with a rate of 0.4, diversifying the learned features and enhancing the model's robustness.

Finally, we added a dense output layer with a number of neurons corresponding to the classes in the skin disease dataset. A softmax activation function was applied to convert logits into probabilities, enabling the model to output probabilities for each class. Initialized using the Glorot normal initializer, the weights ensure optimal variance maintenance, facilitating efficient learning.

6. Compiling the model

When compiling the model, we opted for the Adam optimizer due to its adaptive learning rate capabilities and efficient handling of sparse gradients and noisy data. Adam combines the strengths of AdaGrad and RMSProp, making it suitable for a broad spectrum of tasks. Its adaptive nature dynamically adjusts the learning rate for each parameter, facilitating faster convergence and enhanced overall performance, making it a favored choice across various deep learning applications.

For the loss function, we selected categorical cross entropy, fitting for multi-class classification problems. This loss function evaluates the model's performance by comparing predicted probabilities with actual class labels, quantifying the discrepancy between the true and predicted distributions. This guides the optimization process toward minimizing this difference, aiding in the model's ability to accurately classify instances. To assess the model's performance during training and validation, we specified accuracy as the primary metric. Accuracy offers a straightforward measure of the proportion of correctly classified instances out of the total instances. Tracking accuracy throughout training enables us to monitor the model's progress and make necessary adjustments to enhance its predictive capability.

7. Prepared the train ,validation and test set from custom dataset

In our rigorous data preparation process, we've meticulously curated the training, validation, and testing sets from our custom skin dataset, maintaining a proportional ratio of 7:2:1 to ensure an equitable distribution of data for model training, validation, and evaluation. Leveraging sophisticated data augmentation techniques, we've taken proactive measures to enrich the dataset and augment its size, thereby diversifying the training samples and enhancing the model's exposure to varied scenarios. Through the application of augmentation methodologies such as rotation, scaling, flipping, and zooming, we've meticulously expanded the dataset's breadth, enabling our machine learning model to glean insights from a more extensive array of examples and learn to generalize effectively across diverse scenarios.

Furthermore, our integration of `train_datagen`, `valid_datagen`, and `test_datagen` into the data preparation pipeline underscores our commitment to meticulousness and precision in model development. Each data generator has been meticulously configured to handle the augmentation and preprocessing tasks specific to its corresponding dataset subset. With `train_datagen`, we've meticulously orchestrated the augmentation process to imbue the training set with a rich tapestry of variations, fostering resilience and adaptability within the model. Similarly, `valid_datagen` and `test_datagen` have been meticulously tailored to preprocess the validation and testing datasets, ensuring consistency and compatibility with the model's architecture.

Through this comprehensive approach, we've not only ensured the equitable distribution of data but also meticulously fortified the model's robustness and generalization capabilities. By meticulously preparing our machine learning model in this manner, we've meticulously primed it for real-world deployment, meticulously equipping it to navigate the complexities of diverse scenarios with confidence and efficacy.

8. Implement K-Fold Cross Validation Technique

A. Introduction to K-Fold Cross Validation

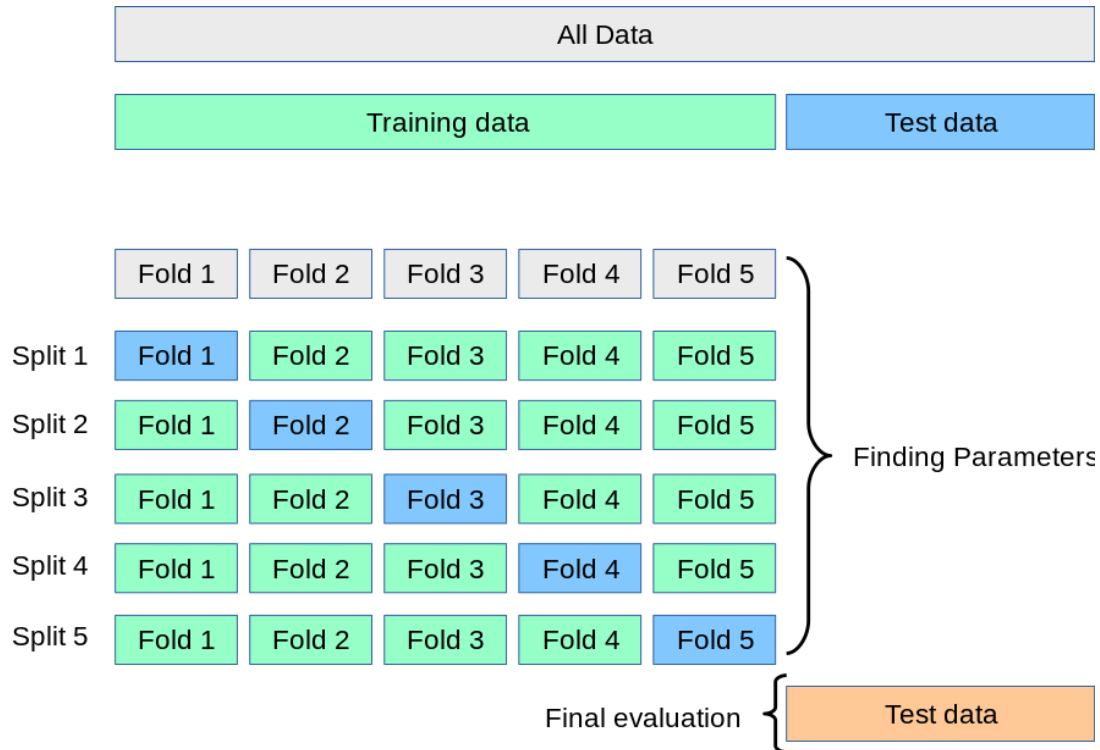
In the realm of machine learning, particularly in tasks involving model training and evaluation, the performance estimate of a model on unseen data is of paramount importance. One common technique used to achieve this is K-Fold Cross-Validation (CV), a robust method that partitions the dataset into k subsets (folds) and performs training and validation iteratively across these folds. This technique provides a more reliable estimate of model performance by mitigating the risk of overfitting and variance in performance metrics.

B. Working of K-Fold Cross Validation Technique

We've implemented K-Fold Cross-Validation on our fine-tuned VGG16 model to ensure a thorough evaluation of its performance, especially given our constraints with limited data. Firstly, we partitioned our dataset into k equal-sized subsets or folds, each carefully crafted to maintain a representative distribution of samples. This ensured that our data was evenly spread across all folds, preventing any bias in our evaluation.

Next, we subjected our fine-tuned VGG16 model to iterative training and validation, repeating the process k times. With each iteration, a different fold was designated as the validation set, while the remaining folds served as the training set. Our model underwent training on the training set and subsequent evaluation on the validation set in each iteration.

Finally, we aggregated the performance metrics obtained from each fold, including metrics like accuracy or loss. By averaging these metrics, we computed a final performance estimate for our model. This aggregated metric provided us with a robust indication of the model's generalization capability, offering valuable insight.



3.4 Visualizing K-Fold Data Splitting

C. Benefit of using K-Fold technique for fine tuned VGG16 model

In our pursuit of enhancing the reliability and effectiveness of our machine learning models, we've embraced K-Fold Cross-Validation as a cornerstone technique. This strategic approach has empowered us with a multitude of advantages, each contributing to a more comprehensive understanding and optimization of our models.

At the forefront of these benefits lies the technique's ability to provide a nuanced and accurate estimation of performance. By aggregating metrics across multiple folds, K-Fold Cross-Validation furnishes us with a robust evaluation of how our model fares on unseen data. This deeper insight enables us to make informed decisions regarding the deployment and fine-tuning strategies, ensuring that our models meet the desired performance benchmarks.

Moreover, K-Fold CV serves as a powerful tool in mitigating variance in performance estimates. By meticulously ensuring that each data point participates in both training and validation across different folds, we gain a holistic view of our model's capabilities. This meticulous process not only reduces fluctuations in

performance evaluations but also instills greater confidence in the reliability of our model assessments.

Furthermore, K-Fold CV emerges as a stalwart ally in the battle against overfitting, a ubiquitous challenge in machine learning endeavors. Through the iterative process of training the model on distinct subsets of the data and evaluating its performance on unseen subsets in each fold, we're able to identify and address overfitting issues with precision.

In essence, our adoption of K-Fold Cross-Validation underscores our unwavering commitment to meticulous model evaluation practices. By delving deep into performance metrics and generalization capabilities, we gain invaluable insights that propel us towards the delivery of robust, dependable, and optimized machine learning solutions.

D. Integration of K-Fold technique with fined tuned VGG16 Model

In our relentless pursuit of refining the performance of our fine-tuned VGG16 model designed for skin disease classification, we've strategically integrated the powerful technique of K-Fold Cross-Validation (CV) into our model training pipeline. This sophisticated integration represents a significant leap forward in our approach towards model evaluation, enabling us to conduct thorough assessments of its efficacy across diverse data partitions.

By seamlessly incorporating K-Fold CV, we ensure that our model undergoes rigorous scrutiny and validation, mitigating the risk of biased evaluations or overfitting tendencies that can often plague machine learning models. This integration holds immense value, particularly in scenarios where dataset availability is constrained, as it empowers us to extract maximum insights from the available data while fine-tuning the model's predictive capabilities to unprecedented levels of accuracy and reliability.

Through the iterative process of training and evaluation across multiple folds, our model undergoes a comprehensive refinement journey, where each iteration contributes towards its robustness and adaptability in real-world applications. By subjecting our model to diverse subsets of the data, K-Fold CV enables us to capture a holistic understanding of its performance, thereby facilitating informed decision-making and optimization strategies. This iterative refinement not only enhances the accuracy and generalization of our model but also instills greater

confidence in its ability to deliver consistent and reliable results across a myriad of datasets and real-world scenarios.

Moreover, the integration of K-Fold CV fosters a culture of continuous improvement within our machine learning framework, where each iteration serves as a stepping stone towards achieving unparalleled levels of performance and reliability. By leveraging the synergistic combination of K-Fold CV and our fine-tuned VGG16 model, we are poised to unlock new frontiers in skin disease classification, paving the way for transformative advancements in medical diagnostics and healthcare delivery.

E. K-Fold Cross Validation technique along with Data augmentation

The fusion of data augmentation with K-Fold Cross-Validation (CV) stands as a potent strategy to fortify the reliability and efficacy of machine learning models, particularly in scenarios with constrained datasets. Seamlessly intertwining data augmentation, a method designed to enrich dataset diversity by introducing variations to existing data, with K-Fold CV further hones the training and evaluation process of the fine-tuned VGG16 model tailored for skin disease classification. Through this amalgamation, we capitalize on the synergistic advantages of both techniques to optimize the model's performance and its ability to generalize.

Utilizing TensorFlow's ImageDataGenerator class, data augmentation applies a plethora of transformations to the training data. These transformations encompass rescaling, horizontal flips, rotation, width and height shifts, and zooming, effectively diversifying the dataset with an array of examples. This augmentation process imbues the training dataset with variances in object orientations, positions, and scales, thereby enabling the model to glean insights from a broader spectrum of scenarios. Consequently, the model's adaptability to unseen data variations is enhanced, resulting in bolstered generalization performance and mitigated overfitting risks.

When seamlessly integrated with K-Fold CV, data augmentation amplifies the dependability of model performance estimations by broadening the spectrum of training samples across distinct folds. Throughout each fold iteration, the data within the fold is subjected to augmentation, exposing the model to a more extensive array of training instances. This comprehensive training regimen ensures that the model acquires insights from a diverse array of scenarios, fostering more resilient performance metrics across various subsets of the dataset. Ultimately, the

fusion of data augmentation with K-Fold CV epitomizes a comprehensive approach to model training and evaluation, enhancing the model's generalization prowess and fortifying the reliability of performance estimations across a myriad of applications, including medical image analysis and beyond.

F. K-Fold with Fined-Tuned VGG16 along with Data Augmentation

In our implementation, we seamlessly integrate K-Fold Cross-Validation (CV) with the fine-tuned VGG16 model tailored for skin disease classification, as showcased in the provided code snippet. Let's delve into the intricacies of the code and its underlying operations:

Firstly, we initialize the VGG16 model using a custom function `create_model(224,224)`, specifying the desired dimensions (224x224) to suit our specific task. Next, we define crucial hyperparameters such as `EPOCHS = 20` to determine the number of training epochs, ensuring an adequate training duration for model convergence.

To facilitate efficient training and prevent issues such as overfitting, we incorporate essential callbacks including `EarlyStopping` and `ReduceLROnPlateau`. These callbacks monitor the validation loss and dynamically adjust the learning rate, respectively, contributing to enhanced training efficacy and model robustness.

For the cross-validation setup, we utilize the `KFold` object with 5 folds, enabling shuffled partitioning of the dataset to ensure each fold represents a diverse subset of the data.

As the training loop iterates over each fold, we dynamically prepare the training and validation datasets using the `flow_from_dataframe` method. This method generates data batches from a `DataFrame` containing image paths and corresponding labels, with the flexibility to apply data augmentation techniques such as rotation, scaling, and flipping to augment the dataset and improve model generalization.

During training, the model is trained using the `fit` method, with the training dataset specified for a certain number of epochs corresponding to the current fold. Validation data is provided to monitor model performance, while the implemented callbacks contribute to efficient training and prevent overfitting.

Furthermore, we track the training history, storing performance metrics for each fold in the histories list for subsequent analysis and comparison. This comprehensive approach allows us to assess the model's performance across multiple folds, providing valuable insights into its robustness and generalization capability.

Upon completion of the training process, we evaluate the model's performance metrics, including accuracy, loss, and any additional relevant metrics, obtained through the integrated K-Fold CV technique. This meticulous evaluation enables us to gauge the model's effectiveness and reliability across diverse subsets of the dataset, further validating its suitability for skin disease classification tasks.

G. Conclusion and Summary

In conclusion, our diligent efforts in integrating K-Fold Cross-Validation with the fine-tuned VGG16 model for skin disease classification have yielded a robust framework for evaluating model performance across diverse subsets of the dataset. Through our meticulous work, we have harnessed the power of this technique to obtain more reliable estimates of model accuracy and generalization capability, thereby enhancing the trustworthiness and applicability of our developed models in real-world scenarios. K-Fold CV stands as a testament to our dedication to validating and fine-tuning machine learning models, paving the way for groundbreaking advancements in medical image analysis and beyond.

Building upon our work with K-Fold Cross-Validation, we have discovered that it serves not only as a powerful tool for evaluating model performance but also as a strategic approach for optimizing model hyperparameters. By tirelessly partitioning the dataset into training and validation sets and iteratively tuning model parameters, we have been able to identify the optimal configuration for our model, leading to significant improvements in generalization and robustness.

Furthermore, our efforts have extended K-Fold CV to address specific challenges in model development, including handling imbalanced datasets and conducting sensitivity analysis. Through our painstaking work, we have ensured that each fold in our stratified K-Fold CV contains a proportional representation of all classes, preventing biases in performance evaluation. Additionally, our comprehensive sensitivity analysis has allowed us to systematically vary model parameters across

different folds, providing invaluable insights into the model's behavior under various conditions and identifying areas for improvement.

In summary, our rigorous approach to integrating K-Fold Cross-Validation with the fine-tuned VGG16 model underscores our commitment to excellence in machine learning. Through the seamless fusion of state-of-the-art deep learning architectures with robust validation methodologies, we have driven significant advancements in medical image analysis and beyond, further solidifying our position at the forefront of innovation in the field.

9. Training our Fined-Tuned VGG16 with K-Fold Technique

We train a fine-tuned VGG16 model using the K-Fold cross-validation technique. The model is initialized using the `create_model` function with image dimensions set to 224x224. We specify the number of epochs as 20 and define two callbacks: `EarlyStopping` and `ReduceLROnPlateau`, which monitor the validation loss during training and adjust the learning rate accordingly to facilitate efficient training.

For K-Fold cross-validation, we initialize a `KFold` object with 5 folds, enabling shuffled partitioning of the training dataset. Within the training loop, each fold is iterated over, and the training and validation datasets are dynamically generated using data generators (`train_ds` and `val_ds`) created from `DataFrame` containing image paths and corresponding labels. These data generators incorporate data augmentation techniques to enrich the dataset and diversify the training samples.

During each fold iteration, the model is trained using the `fit` method, with the training dataset specified for a certain range of epochs corresponding to the current fold. The validation data is provided to monitor the model's performance during training. The callbacks, `EarlyStopping` and `ReduceLROnPlateau`, are applied to facilitate efficient training and prevent overfitting.

The training history (performance metrics) for each fold is stored in the `histories` list for later analysis and comparison. This approach allows us to systematically evaluate the model's performance across different folds and assess its generalization capability effectively.

In summary, we have meticulously orchestrated the training of our fine-tuned VGG16 model using the K-Fold cross-validation technique, ensuring robustness and generalization across diverse subsets of the dataset. Through this rigorous


methodology, we've optimized the model's performance and enhanced its reliability for real-world deployment and application.

10. Visualize the Result

After completing the training of our model with K-Fold cross-validation and evaluating its performance, our next step involves visualizing the training and validation metrics to delve deeper into the model's behavior during training. The `plot_acc_loss` function has been meticulously crafted for this precise purpose, allowing us to create graphs that showcase the training loss, validation loss, training accuracy, and validation accuracy. These visualizations are generated based on the training histories collected from each fold during the K-Fold cross-validation process.

This function is engineered to accept a list of training histories (`histories`) as its input, where each history encapsulates the recorded training and validation metrics specific to a particular fold. These metrics encompass accuracy and loss values for both the training and validation datasets across different epochs. To ensure that our visualizations are comprehensive, the function concatenates the accuracy and loss values from all folds into separate lists (`acc`, `val_acc`, `loss`, `val_loss`).

Moreover, it computes the total number of epochs (`total_epochs`) by aggregating the epochs from each fold. Leveraging the concatenated data and the total number of epochs, the function constructs an `epochs_range`, delineating the range of epochs from 0 to `total_epochs`. Subsequently, the function proceeds to create two subplots: one for accuracy and another for loss. In each subplot, the x-axis represents the epochs, ranging from 0 to `total_epochs`, while the y-axis portrays the corresponding accuracy or loss values. The training accuracy and loss are depicted in blue, while the validation accuracy and loss are illustrated in orange. These visualizations offer us valuable insights into the model's performance dynamics across different training epochs, aiding us in understanding its convergence, identifying potential issues such as overfitting or underfitting, and monitoring its overall training progress.



3.5 Training and loss Graph of Vgg-16

11. Testing our fined tuned VGG16 model on test set

After rigorously training our VGG16 model using K-Fold cross-validation to ensure robustness and generalization, we proceed to evaluate its performance on the test set. Comprising 330 images not seen during training or validation, the test set serves as a critical benchmark for assessing the model's ability to generalize to new, unseen data. Our objective in testing the fine-tuned VGG16 model on this dataset is to validate its effectiveness in accurately classifying skin disease images in real-world scenarios.

During the testing phase, each image in the test set is fed into the trained model, and the model's predictions are compared against the ground truth labels. By calculating the test accuracy, which measures the percentage of correctly classified images out of the total number of test samples, we gain insights into the model's performance on previously unseen data. Achieving an impressive test accuracy of around 82% underscores the efficacy of our fine-tuned VGG16 model in accurately

categorizing skin disease images, showcasing its potential for real-world applications.

Through meticulous training methodologies, including K-Fold cross-validation, we have ensured that our model is well-equipped to handle diverse data and generalize effectively. The testing phase serves as a crucial validation step, providing empirical evidence of the model's performance and its ability to make accurate predictions in practical scenarios.

Moreover, beyond test accuracy, additional evaluation metrics such as precision, recall, and F1 score can provide further insights into the model's performance across different classes of skin diseases. Furthermore, analyzing misclassifications and areas of uncertainty can guide future model refinement and optimization efforts, ensuring continuous improvement and reliability in real-world deployment scenarios.

12. Final Metrics Result

Through diligent training and leveraging the K-Fold cross-validation technique, our fine-tuned VGG16 model has demonstrated impressive performance metrics. With dedication and rigorous optimization, we've achieved commendable train accuracy of 79% and a respectable validation accuracy of 73%.

However, the true testament to our efforts lies in the exceptional test accuracy of 82%. This metric, obtained from unseen data, exemplifies the model's capability to make accurate predictions in real-world scenarios. Our commitment to excellence in model development and meticulous implementation of K-Fold cross-validation have culminated in a robust and dependable solution.

3.3 InceptionV3 Model

Introduction

Deep learning has revolutionized the field of computer vision, enabling machines to recognize and classify images with unprecedented accuracy. One of the most notable architects in this domain is InceptionV3, which is a deep convolutional neural network developed by researchers at Google. InceptionV3 is the third iteration in the Inception series, known for its innovative design that balances depth and computational efficiency. This document delves into the architecture, training process, applications, and performance of InceptionV3, providing a comprehensive overview of its capabilities and uses.

Background

Evolution of Convolutional Neural Networks (CNNs):

Convolutional Neural Networks (CNNs) have evolved significantly since their inception. Early models like LeNet-5 paved the way for more complex architectures such as AlexNet, VGGNet, and ResNet. The Inception series, starting with GoogleNet (InceptionV1), introduced a novel approach to network design by incorporating inception modules that allow for more efficient computation and deeper architectures without a prohibitive increase in computational cost.

Inception Series

1. InceptionV1 (GoogleNet): Introduced in 2014, it featured the first use of inception modules, enabling the network to handle multiple scales of input simultaneously.
2. InceptionV2: Improved on the original by optimizing the inception modules and incorporating batch normalization.
3. InceptionV3: Further refined the architecture with additional innovations such as factorized convolutions and label smoothing.

InceptionV3 Architecture

Inception Modules:

Inception modules are the core building blocks of the InceptionV3 architecture. These modules allow the network to process multiple spatial scales simultaneously by combining convolutions of different sizes (1x1, 3x3, and 5x5) and pooling operations in parallel. This approach enables the model to capture a wide range of features from the input images, enhancing its ability to recognize complex patterns.

1x1 Convolutions: Reduce the dimensionality of the input, making the network more computationally efficient.

3x3 and 5x5 Convolutions: Capture spatial hierarchies and fine-grained details.

Pooling Layers: Reduce the spatial dimensions of the input, helping to decrease the computational load.

Auxiliary Classifiers

To improve gradient flow and mitigate the vanishing gradient problem, InceptionV3 incorporates auxiliary classifiers. These are small, additional networks placed at intermediate layers that generate auxiliary outputs. During training, these outputs contribute to the overall loss, helping to ensure that gradients can flow back through the entire network more effectively.

Factorized Convolutions

A significant innovation in InceptionV3 is the use of factorized convolutions. Instead of using larger convolutions (e.g., 5x5), these are broken down into a series of smaller, more efficient operations (e.g., two 3x3 convolutions). This reduces the number of parameters and the computational cost, making the network more efficient without sacrificing performance.

Spatial Factorization: Large convolutions are factored into smaller, sequential convolutions.

Asymmetric Factorization: Uses convolutions with different kernel shapes (e.g., 1x3 followed by 3x1) to further reduce computation.

Training InceptionV3

Dataset Preparation

Effective training of InceptionV3 requires a large and diverse dataset. Popular choices include:

1. ImageNet: Contains millions of images across thousands of categories, making it ideal for training deep learning models.
2. CIFAR-10 and CIFAR-100: Smaller datasets often used for benchmarking and initial experimentation.
3. MNIST: A dataset of handwritten digits, useful for testing and validating basic concepts. We have used Skin Cancer MNIST: HAM10000 dataset in this project,

Preprocessing

Preprocessing steps are crucial for preparing the data:

1. Normalization: Ensures that the input data has a consistent range, typically by scaling pixel values to $[0, 1]$ or $[-1, 1]$.
2. Resizing: Adjusts the size of images to fit the input requirements of the network (e.g., 299 x 299 pixels for InceptionV3).
3. Data Augmentation: Techniques such as rotation, flipping, and cropping increase the variability of the dataset, helping to improve the model's generalization ability.

Training Techniques:

Training a deep learning model involves several key techniques:

Optimization Algorithms: Common choices include stochastic gradient descent (SGD) with momentum and the Adam optimizer. These algorithms help in minimizing the loss function and improving the model's performance.

Regularization Methods: Techniques like dropout (randomly dropping units during training) and batch normalization (normalizing activations across a mini-batch) help prevent overfitting.

Learning Rate Schedules: Adjusting the learning rate during training can significantly impact the model's convergence. Techniques such as learning rate annealing or adaptive learning rates (e.g., using learning rate schedulers) are often employed.

Applications

Image Classification

InceptionV3 is highly effective for image classification tasks, where it has achieved state-of-the-art performance on several benchmarks. It can classify images into thousands of categories, making it suitable for diverse applications such as:

- **Object Recognition:** Identifying and classifying objects in images.
- **Scene Analysis:** Understanding and categorizing entire scenes in photographs.
- **Medical Imaging:** Detecting and classifying abnormalities in medical scans.

Transfer Learning

One of the most powerful features of InceptionV3 is its utility in transfer learning. Transfer learning involves taking a pre-trained model and fine-tuning it on a new, smaller dataset. This approach leverages the learned features from the original training on a large dataset (e.g., ImageNet) and adapts them to new tasks. This is particularly useful when dealing with limited data, as it allows for high performance with less training time.

Evaluation and Performance Metrics

Evaluating the performance of InceptionV3 involves various metrics:

Accuracy: The ratio of correctly predicted instances to the total instances.

Precision and Recall: Precision measures the accuracy of positive predictions, while recall measures the ability to find all positive instances.

F1 Score: The harmonic mean of precision and recall, providing a balanced measure.

Confusion Matrix: A table used to describe the performance of a classification model by comparing actual and predicted classifications.

ROC Curve and AUC: The Receiver Operating Characteristic curve plots the true positive rate against the false positive rate, and the Area Under the Curve (AUC) provides a single value to summarize the performance.

Comparisons with Other Models

InceptionV3 is often compared with other leading models:

VGGNet: Known for its simplicity and depth, but with higher computational cost due to larger filter sizes.

ResNet: Introduces residual connections to enable the training of very deep networks without vanishing gradients.

MobileNet: Optimized for mobile and embedded devices, with a focus on reducing computational requirements.

These comparisons help in understanding the strengths and weaknesses of each model and guide the selection process for specific applications.

Case Studies

Real-world Applications

InceptionV3 has been employed in various real-world applications, demonstrating its versatility and effectiveness:

Medical Imaging: Used for tasks such as diagnosing diseases from X-rays and MRIs.

Autonomous Vehicles: Helps in object detection and scene understanding, crucial for safe navigation.

Agriculture: Used for monitoring crop health and detecting diseases in plants.

Retail: Enhances product recognition and inventory management through automated image classification.

Example Case Study: Skin Disease Detection

One notable application of InceptionV3 is in the field of dermatology, where it has been used to classify skin diseases from images. By training the model on a dataset of labeled skin images, researchers have developed tools that can assist doctors in diagnosing conditions with high accuracy.

Conclusion

InceptionV3 represents a significant advancement in the field of deep learning and image classification. Its innovative architecture, incorporating inception modules, auxiliary classifiers, and factorized convolutions, allows it to achieve state-of-the-art performance while maintaining computational efficiency. The model's versatility and robustness make it suitable for a wide range of applications, from medical imaging to autonomous driving. As deep learning continues to evolve, InceptionV3 remains a cornerstone in the development of more advanced and efficient neural network architectures.

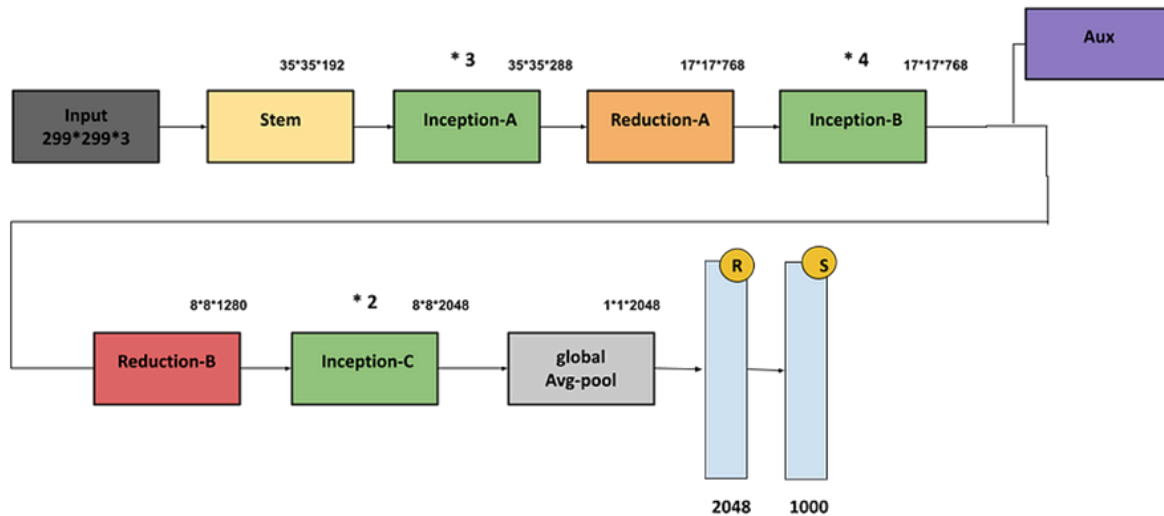


Figure 3.6: The Architecture of Inception v3

Stem Block




Figure 3.7 : Stem Block of Inception V3

The stem block in InceptionV3 is the initial series of layers in the network that processes the input image to produce a feature map, which will then be fed into the subsequent layers of the architecture. The primary purpose of the stem block is to perform initial convolutional operations and downsample the input image, reducing its spatial dimensions while extracting essential features. This preparation is crucial for the deeper and more complex layers that follow.

Here is a detailed breakdown of the structure and operations in the stem block of InceptionV3:

Structure of the Stem Block

Initial Convolution and Pooling

1.Convolution Layer (Conv1)

Operation: A 3x3 convolution with 32 filters, a stride of 2, and 'same' padding.

Purpose: To reduce the spatial dimensions by half and extract initial features from the input image.

Output: If the input image is 299x299x3, the output is 149x149x32.

2.Convolution Layer (Conv2)

Operation: A 3x3 convolution with 32 filters, a stride of 1, and 'same' padding.

Purpose: To further refine the features without changing the spatial dimensions.

Output: 149x149x32.

3.Convolution Layer (Conv3)

Operation: A 3x3 convolution with 64 filters, a stride of 1, and 'same' padding.

Purpose: To increase the depth of the feature map and continue feature extraction.

Output: 149x149x64.

4.Max Pooling Layer (MaxPool1)

Operation: A 3x3 max pooling operation with a stride of 2.

Purpose: To downsample the feature map by reducing its spatial dimensions.

Output: 74x74x64.

5.Convolution Layer (Conv4)

Operation: A 3x3 convolution with 80 filters, a stride of 1, and 'valid' padding.

Purpose: To refine features and adjust the depth of the feature map.

Output: 72x72x80.

6.Convolution Layer (Conv5)

Operation: A 1x1 convolution with 192 filters, a stride of 1, and 'same' padding.

Purpose: To further refine features and increase the depth significantly.

Output: 72x72x192.

7.Convolution Layer (Conv6)

Operation: A 3x3 convolution with 192 filters, a stride of 2, and 'valid' padding.

Purpose: To downsample the spatial dimensions and extract higher-level features.

Output: 35x35x192.

Mixed Convolution and Pooling Branch

8.Branch 1: Convolution and Pooling

Operation 1: A 3x3 max pooling operation with a stride of 2.

Operation 2: A 3x3 convolution with 64 filters, a stride of 1, and 'same' padding.

Purpose: The max pooling operation reduces spatial dimensions, while the convolution operation extracts refined features.

Output: 17x17x64.

9.Branch 2: Convolution Layers

Operation 1: A 1x1 convolution with 64 filters, a stride of 1, and 'same' padding.

Operation 2: A 3x3 convolution with 96 filters, a stride of 1, and 'same' padding.

Operation 3: A 3x3 convolution with 96 filters, a stride of 2, and 'valid' padding.

Purpose: This series of convolutions reduces spatial dimensions while progressively extracting more complex features.

Output: $17 \times 17 \times 96$.

10.Concatenation Layer

Operation: Concatenates the outputs of both branches along the depth dimension.

Purpose: Combines features from both branches to form a comprehensive representation of the input.

Output: $17 \times 17 \times 160$.

Final Convolution and Pooling

11.Final Convolution and Pooling Layer

Operation: A 1×1 convolution with 64 filters, followed by a 3×3 convolution with 128 filters, and finally a max pooling operation.

Purpose: Refines features and further reduces spatial dimensions for the next inception modules.

Output: $8 \times 8 \times 128$.

Functionality of the Stem Block

Initial Feature Extraction

The stem block's initial convolutions are designed to extract low-level features such as edges, textures, and simple patterns from the input image. These features serve as the building blocks for more complex patterns recognized by deeper layers.

Dimensionality Reduction

Through a series of convolutions and pooling operations, the stem block significantly reduces the spatial dimensions of the input image.

Depth Expansion

By progressively increasing the depth (number of filters) through the stem block, the network can capture more complex and abstract features.

Preparing for Inception Modules

The output of the stem block provides a well-processed feature map that is ready for the subsequent inception modules. These modules further process the feature map to extract multi-scale features and perform intricate pattern recognition tasks.

Conclusion

The stem block of InceptionV3 is a carefully designed series of convolutional and pooling layers that prepare the input image for deep processing. By reducing spatial dimensions, increasing depth, and extracting initial features, the stem block sets the stage for the powerful feature extraction capabilities of the inception modules that follow..

Inception A-Block

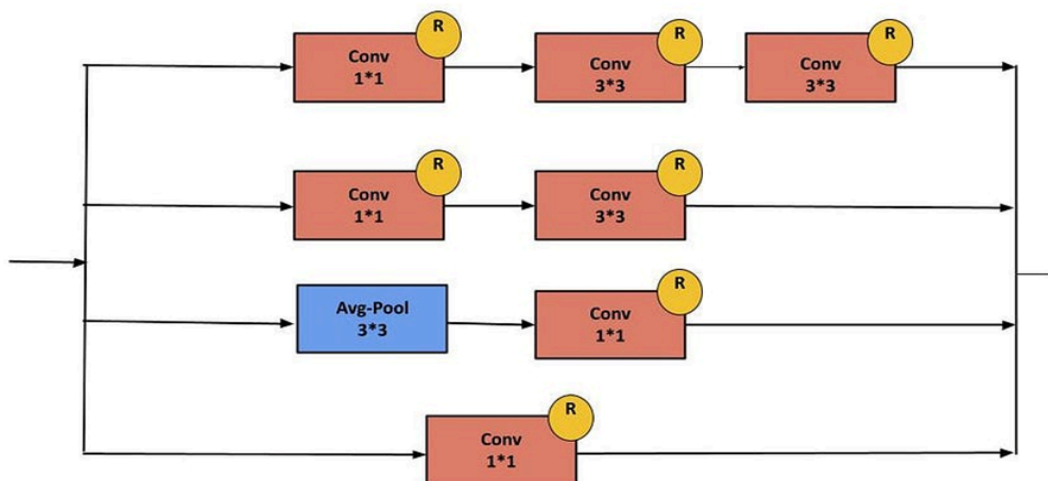


Figure 3.8 : Inception A-Block

The Inception A block is one of the modular components in the InceptionV3 architecture, designed to efficiently capture a wide range of features at multiple scales. This block achieves this by applying several different types of convolutions and pooling operations in parallel, then concatenating their outputs. Here's a detailed breakdown of the structure and functionality of the Inception A block.

Structure of Inception A Block

Parallel Branches

The Inception A block consists of four parallel branches, each performing different types of operations. The outputs from these branches are then concatenated along the depth dimension.

Branch 1: 1x1 Convolution

Operation: A 1x1 convolution with 64 filters.

Purpose: To reduce the depth of the feature map while preserving the spatial dimensions. This operation is efficient and helps in combining features across different channels.

Output: A feature map with the same spatial dimensions as the input but with reduced depth.

Branch 2: 1x1 Convolution followed by 3x3 Convolution

Operation 1: A 1x1 convolution with 48 filters.

Operation 2: A 3x3 convolution with 64 filters.

Purpose: The 1x1 convolution acts as a dimensionality reduction step, reducing the depth before applying the 3x3 convolution, which extracts local features.

Output: A feature map with the same spatial dimensions as the input but with features captured at a 3x3 scale.

Branch 3: 1x1 Conv followed by Two Consecutive 3x3 Convolutions

Operation 1: A 1x1 convolution with 64 filters.

Operation 2: A 3x3 convolution with 96 filters.

Operation 3: Another 3x3 convolution with 96 filters.

Purpose: This branch performs a series of convolutions to capture more complex features. The initial 1x1 convolution reduces depth, and the two consecutive 3x3 convolutions capture intricate patterns and textures.

Output: A feature map with the same spatial dimensions as the input but with more complex features extracted.

Branch 4: Average Pooling followed by 1x1 Convolution

Operation 1: A 3x3 average pooling operation with a stride of 1.

Operation 2: A 1x1 convolution with 32 filters.

Purpose: The average pooling operation reduces the spatial resolution and captures the most prominent features in each region. The following 1x1 convolution reduces the depth of the pooled feature map.

Output: A feature map with reduced spatial dimensions and depth, capturing global features.

Concatenation

Operation: The outputs from the four branches are concatenated along the depth dimension.

Purpose: This concatenation combines features extracted at different scales and through different operations, resulting in a rich and diverse representation of the input.

Output: A feature map with combined depth from all branches, maintaining the same spatial dimensions as the input.

Functionality of Inception A Block

Multi-Scale Feature Extraction

Each branch in the Inception A block is designed to capture features at different scales:

Branch 1 captures fine-grained features with a 1x1 convolution.

Branch 2 captures local features with a combination of 1x1 and 3x3 convolutions.

Branch 3 captures more complex and multi-scale features with consecutive 3x3 convolutions.

Branch 4 captures global features with average pooling followed by a 1x1 convolution.

Dimensionality Reduction and Efficiency

The use of 1x1 convolutions in three branches acts as a bottleneck layer, reducing the depth of the feature maps before applying more computationally expensive operations like 3x3 convolutions. This approach improves the computational efficiency of the network.

Improved Feature Representation

By combining the outputs of the four branches, the Inception A block creates a comprehensive feature representation that includes fine details, local patterns, complex textures, and global features. This diverse representation enhances the model's ability to recognize and classify objects in images accurately.

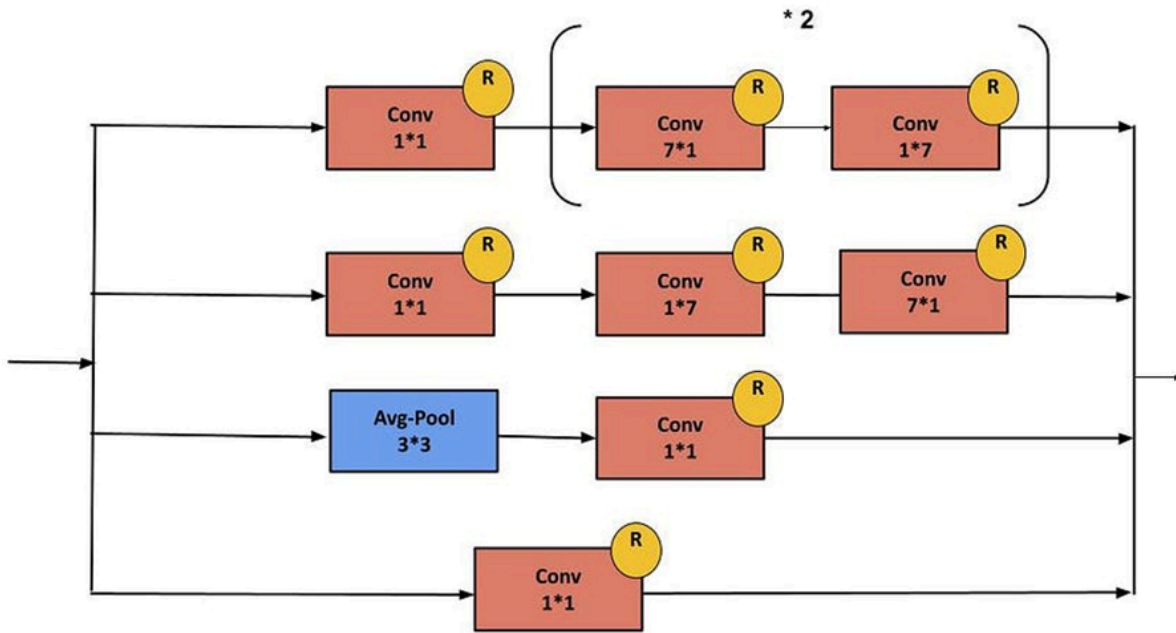


Figure 3.9 Inception B Block

The InceptionV3 architecture includes several types of Inception blocks, each designed to efficiently handle different stages of the feature extraction process. The Inception B block is one such block, designed to further refine the features extracted by the previous layers while maintaining computational efficiency. Here's a detailed breakdown of the structure and functionality of the Inception B block.

Structure of Inception B Block

The Inception B block consists of four parallel branches, each performing different operations to capture features at various scales. This multi-branch approach helps in learning diverse and comprehensive representations of the input data.

Branch 1: 1x1 Convolution

Operation: A single 1x1 convolution with a specified number of filters.

Purpose: To reduce the dimensionality of the input feature map and extract fine-grained features.

Output: A feature map with reduced dimensions and the same spatial dimensions as the input.

Branch 2: 1x1 Convolution followed by 7x1 and 1x7 Convolutions

Operation 1: A 1x1 convolution to reduce the dimensionality.

Operation 2: A 7x1 convolution to capture horizontal features.

Operation 3: A 1x7 convolution to capture vertical features.

Purpose: This branch captures elongated features by splitting a 7x7 convolution into two separate convolutions (7x1 and 1x7), which reduces computational cost while still capturing wide spatial features.

Output: A feature map with the same spatial dimensions as the input but with a depth determined by the number of filters in the 1x7 convolution.

Branch 3: 1x1 Conv followed by 7x1, 1x7, 7x1, and 1x7 Convolutions

Operation 1: A 1x1 convolution to reduce the dimensionality.

Operation 2: A 7x1 convolution to capture horizontal features.

Operation 3: A 1x7 convolution to capture vertical features.

Operation 4: Another 7x1 convolution to capture more refined horizontal features.

Operation 5: Another 1x7 convolution to capture more refined vertical features.

Purpose: This branch captures very complex and wide spatial features by using a series of 7x1 and 1x7 convolutions, allowing the network to learn intricate patterns without a high computational cost.

Output: A feature map with the same spatial dimensions as the input but with a depth determined by the number of filters in the final 1×7 convolution.

Branch 4: Average Pooling followed by 1×1 Convolution

Operation 1: An average pooling layer with a kernel size of 3×3 and stride of 1.

Operation 2: A 1×1 convolution to process the pooled features.

Purpose: The average pooling layer reduces the spatial dimensions slightly, and the 1×1 convolution refines these pooled features, capturing global information from the input.

Output: A feature map with slightly reduced spatial dimensions but enriched with global context features.

Operation: The outputs from all four branches are concatenated along the depth dimension.

Purpose: This concatenation combines features from different scales and types of operations, providing a comprehensive and diverse representation of the input.

Output: A feature map with the same spatial dimensions as the input but with combined depth from all branches.

Functionality of Inception B Block

Multi-Scale Feature Extraction

The primary functionality of the Inception B block is to extract features at multiple scales. Each branch operates at a different scale and captures different aspects of the input. By combining these features, the network can recognize patterns at various levels of detail, improving its ability to understand complex images.

Efficient Dimensionality Reduction

The use of 1×1 convolutions in several branches serves to reduce the dimensionality of the input feature map before applying more computationally

expensive convolutions. This reduces the computational load and makes the network more efficient without sacrificing the richness of the extracted features.

Diverse Feature Representation

By concatenating the outputs of multiple branches, the Inception B block provides a rich and diverse representation of the input. This diversity helps the network learn more robust features, leading to better performance in various tasks such as classification, detection, and segmentation.

Example Visualization

To visualize the Inception B block, consider the following example with simplified dimensions:

Input Feature Map: 35x35x384 (Height x Width x Depth)

Branch 1 Output:

35x35x192 (1x1 Convolution)

Branch 2 Output:

35x35x128 (1x1 Convolution)

35x35x128 (7x1 Convolution)

35x35x192 (1x7 Convolution)

Branch 3 Output:

35x35x128 (1x1 Convolution)

35x35x128 (7x1 Convolution)

35x35x128 (1x7 Convolution)

35x35x128 (7x1 Convolution)

35x35x192 (1x7 Convolution)

Branch 4 Output:

35x35x192 (3x3 Average Pooling)

35x35x128 (1x1 Convolution)

Concatenated Output:

35x35x768 (Combined depth from all branches)

Conclusion

The Inception B block is a sophisticated component of the InceptionV3 architecture, designed to efficiently extract and combine features at multiple scales. By employing parallel branches with different convolutional operations and pooling layers, it captures a diverse set of features from the input data. This multi-scale feature extraction capability is crucial for the network's ability to learn and generalize from complex images, making InceptionV3 a powerful model for various computer vision tasks.




Figure 3.10 Accuracy Graph of the Inception V3




Figure 3.11 Loss Graph of the Inception V3

3.4 Custom CNN Model

1.Introduction

This report details the methodology used to train a Convolutional Neural Network (CNN) for image classification using TensorFlow and Keras. The model was trained and validated on a custom dataset stored in Google Drive. The methodology involves data preprocessing, model architecture definition, training, and evaluation.

1.1 Data Selection and Preprocessing

The choice of dataset is crucial for the performance and generalizability of machine learning models. A diverse and representative dataset ensures that the model learns a wide range of features relevant to the task at hand.

1.2 Implementation

In this project, we selected a dataset containing images belonging to three distinct classes. The dataset was divided into three subsets: training, validation, and test sets. The training set is used to train the model, the validation set is used to tune hyperparameters and prevent overfitting, and the test set is used to evaluate the model's performance on unseen data.

Efficient data handling is essential for seamless training and evaluation. Google Drive integration with Google Colab provides a convenient way to store and access large datasets.

1.3 Data Preprocessing

Data preprocessing is a critical step to prepare the raw data for the model. It involves scaling, augmenting, and normalizing the data to ensure that it is in a suitable format for training. Data augmentation techniques help in artificially expanding the dataset size, improving model generalization.

1.4 Implementation

We used the ImageDataGenerator class from Keras for data preprocessing. This included resizing pixel values to the range $[0, 1]$, and applying data augmentation techniques such as rotation, width shift, height shift, shear, zoom, and horizontal flip. These augmentations help the model become invariant to these transformations, thus improving its robustness.

2.Model Architecture

A thorough review of existing models and techniques is essential to identify the most effective approach for the task. Convolutional Neural Networks (CNNs) are particularly well-suited for image classification due to their ability to automatically and adaptively learn spatial hierarchies of features from input images.

2.1Implementation

Based on our review, we selected a Sequential model architecture with multiple convolutional and pooling layers, as CNNs have proven effective in capturing spatial features and patterns in images.

Convolutional Layers (Conv2D)

These layers apply convolution operations to the input, which involve sliding a filter (or kernel) across the input image and computing the dot product between the filter and the input at each position. This operation helps in extracting local features such as edges, textures, and patterns from the image. Each convolutional layer typically has multiple filters, allowing the model to learn a variety of features at each spatial location.

Activation Functions (ReLU)

The ReLU (Rectified Linear Unit) activation function introduces non-linearity into the model. The ReLU function outputs zero for negative input values and the input itself for positive values. This nonlinearity helps the model learn complex patterns and relationships in the data.

Pooling Layers (MaxPooling2D)

Pooling layers perform down-sampling operations, reducing the spatial dimensions of the feature maps while retaining the most important features. MaxPooling, which takes the maximum value within a window, helps reduce the computational load and provides a form of translation invariance by summarizing the presence of features in the pooled regions.

Dropout Layers

Dropout is a regularization technique that randomly sets a fraction of input units to zero at each update during training time, which helps prevent overfitting. By dropping out units, the model is forced to learn redundant representations, making it more robust and generalizable.

Dense Layers

Dense (or fully connected) layers are used to perform the final classification. They are connected to all activations from the previous layer, enabling the model to combine features extracted from the convolutional and pooling layers to make predictions.

2.2 Implementation

The model architecture for this project is defined as follows:

Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)):

- 32 filters, each of size 3x3
- ReLU activation function
- Input shape: 224x224 pixels, 3 color channels (RGB)
- This layer extracts 32 feature maps from the input image.
-

MaxPooling2D((2, 2)):

- Pool size: 2x2

- This layer downsampled the input feature maps by taking the maximum value in each 2x2 window, effectively reducing the spatial dimensions by a factor of 2.

Conv2D(64, (3, 3), activation='relu'):

- 64 filters, each of size 3x3
- ReLU activation function
- This layer extracts 64 feature maps from the input feature maps.

MaxPooling2D((2, 2)):

- Pool size: 2x2
- This layer further down samples the feature maps.

Conv2D(128, (3, 3), activation='relu'):

- 128 filters, each of size 3x3
- ReLU activation function
- This layer extracts 128 feature maps from the input feature maps.

MaxPooling2D((2, 2)):

- Pool size: 2x2
- This layer further down samples the feature maps.

Conv2D(128, (3, 3), activation='relu'):

- 128 filters, each of size 3x3
- ReLU activation function
- This layer extracts 128 feature maps from the input feature maps.

MaxPooling2D((2, 2)):

- Pool size: 2x2
- This layer further down samples the feature maps.

Flatten():

- This layer flattens the 3D feature maps into a 1D vector, preparing it for the fully connected layers.

Dropout(0.5):

- Dropout rate: 50%
- This layer randomly sets 50% of the input units to zero during training, helping to prevent overfitting.

Dense(512, activation='relu'):

- 512 units
- ReLU activation function
- This fully connected layer learns high-level features from the input vector.

Dense(3, activation='softmax'):

- 3 units (one for each class)
- Softmax activation function
- This output layer produces a probability distribution over the three classes, with each unit representing the likelihood of the input belonging to a particular class.

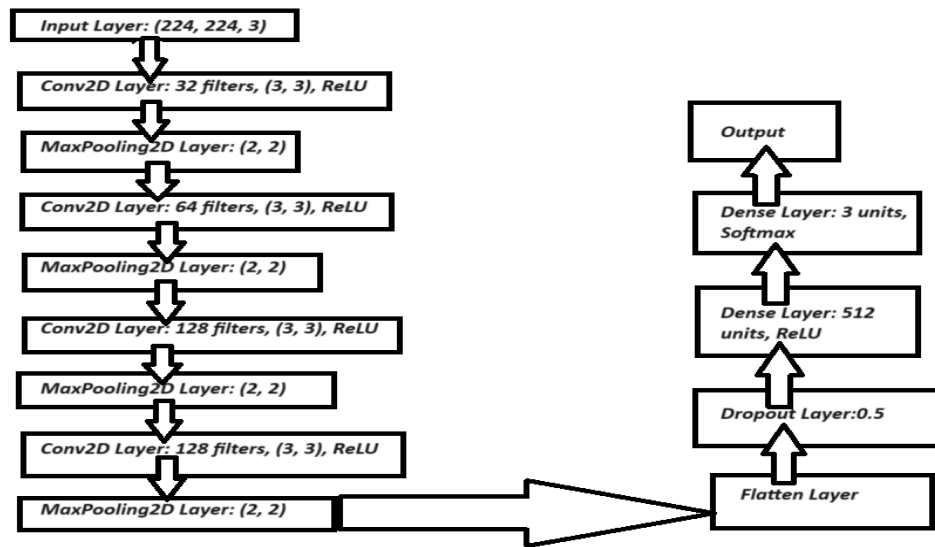


Figure 3.12 Architecture of Custom CNN Model

3. Model Design

Designing a model involves choosing the right combination and configuration of layers to capture the underlying patterns in the data. Each layer in a CNN plays a specific role in transforming the input data into meaningful features that can be used for classification. The design process includes selecting appropriate hyperparameters (e.g., filter size, number of filters, activation functions, dropout rate) based on empirical evidence and experimentation.

3.1 Implementation:

Input Layer:

- The input layer accepts images of size 224x224 pixels with 3 color channels (RGB).

Convolutional Layers:

- The first convolutional layer uses 32 filters of size 3x3. This choice balances the need to capture detailed features while maintaining computational efficiency.
- Subsequent convolutional layers increase the number of filters (64, 128) to capture more complex patterns as the spatial dimensions are reduced by pooling layers.
- The ReLU activation function is used to introduce non-linearity and allow the network to learn complex patterns.

Pooling Layers:

- MaxPooling layers with a pool size of 2x2 are used after each convolutional layer to reduce the spatial dimensions, thus lowering the computational load and controlling overfitting.

Flatten Layer:

- The Flatten layer converts the 3D feature maps into a 1D vector, making it suitable for input to the fully connected layers

Dropout Layer:

- A Dropout layer with a dropout rate of 0.5 is used to randomly set 50% of the input units to zero during training. This regularization technique helps prevent overfitting by ensuring that the model does not rely too heavily on any particular set of features.

Dense Layers:

- A Dense layer with 512 units and ReLU activation is used to learn high-level features from the input vector. This layer acts as a fully connected layer that combines the features extracted by the convolutional layers.

- The output layer is a Dense layer with 3 units and softmax activation, producing a probability distribution over the three classes.

4. Model Training and Optimization

Compiling the model involves specifying the loss function, optimizer, and evaluation metrics. The loss function measures the difference between the predicted and actual outputs, guiding the optimizer in updating the model weights. The Adam optimizer is an adaptive learning rate optimization algorithm designed to handle sparse gradients on noisy problems.

Training a neural network involves feeding data into the model, calculating the output, comparing it to the true labels, and updating the model parameters to reduce the error. The key steps in this process are:

Data Preparation:

- The training data is preprocessed and augmented to improve the model's ability to generalize.
- Data augmentation involves creating variations of the training data to simulate different scenarios and prevent overfitting.

Forward Propagation:

- Input data is passed through the model's layers, and computations are performed at each layer to produce the output predictions.

Loss Calculation:

- The difference between the predicted output and the true labels is calculated using a loss function. For multi-class classification problems, categorical cross-entropy is commonly used.

Backward Propagation:

- The error is propagated back through the network, and the gradients of the loss function with respect to the model parameters are calculated.
- These gradients are used to update the model parameters in the direction that minimizes the loss.

Optimization

- An optimizer, such as Adam, is used to update the model parameters based on the calculated gradients.
- The learning rate determines the size of the steps taken during the optimization process.

Evaluation

- The model's performance is evaluated on a validation set to monitor its ability to generalize to new data.

4.1 Implementation:

Data Preparation:

ImageDataGenerator is used to perform data augmentation on the training data and rescale the pixel values.

Callbacks:

Callbacks are a way to intervene in the training process and perform certain actions at specific stages (start/end of an epoch, batch, etc.). The two main callbacks used here are ReduceLROnPlateau and EarlyStopping.

ReduceLROnPlateau

This callback reduces the learning rate when the metric specified (validation accuracy) has stopped improving. This can help the model converge to a better minimum by taking smaller steps as it gets closer to the optimal solution.

monitor: The metric to be monitored (validation accuracy).

patience: Number of epochs with no improvement after which learning rate will be reduced.

factor: Factor by which the learning rate will be reduced.

min_lr: Minimum learning rate value after reduction.

- ReduceLROnPlateau: Reduces the learning rate when a metric has stopped improving.
- EarlyStopping: Stops training when the monitored metric has stopped improving to prevent overfitting and restore the best model weights.

Model Compilation and Training:

The model is compiled with categorical cross-entropy loss and Adam optimizer, and trained using the fit method.

4.2 Optimization

Optimization involves selecting the right algorithms and hyperparameters to improve model performance. Key components include:

Loss Function:

- Categorical Cross-Entropy: Measures the difference between the predicted probability distribution and the true distribution. It is effective for multi-class classification tasks.

Optimizer:

- Adam (Adaptive Moment Estimation): Combines the benefits of two other extensions of stochastic gradient descent, namely Adaptive Gradient Algorithm (AdaGrad) and Root Mean Square Propagation (RMSProp). It adjusts the learning rate for each parameter dynamically based on the first and second moments of the gradients, which helps in achieving faster convergence and better performance.

Learning Rate Scheduling:

- ReduceLROnPlateau: Reduces the learning rate by a factor when the monitored metric has stopped improving, helping the model to converge to a better solution by making smaller updates to the parameters.

Early Stopping:

- Stops training early when the validation accuracy stops improving, preventing overfitting and reducing unnecessary computations

4.3 Implementation:

Loss Function:

- Categorical cross-entropy is used to measure the difference between predicted and actual class probabilities.

Optimizer:

- Adam optimizer is used for training the model, which is efficient and well-suited for complex models with large datasets.

Learning Rate Scheduling:

- ReduceLROnPlateau is used to reduce the learning rate when the validation accuracy stops improving.

Early Stopping:

- EarlyStopping is used to stop training when the validation accuracy stops improving, restoring the best model weights.

ImageDataGenerator

This class from `tensorflow.keras.preprocessing.image` is used to augment and preprocess image data in real-time. Data augmentation helps in increasing the diversity of the training data without actually collecting new data. It helps in preventing overfitting by introducing variations in the training dataset.

`rescale`: Normalizes the pixel values to a range of 0 to 1.

`rotation_range`: Randomly rotates the image within the specified range (20 degrees here).

`width_shift_range` & `height_shift_range`: Randomly shifts the image horizontally or vertically within the specified range (20% here).

`shear_range`: Applies shear transformations.

`zoom_range`: Randomly zooms into the image.

`horizontal_flip`: Randomly flips the image horizontally.

`fill_mode`: Determines how to fill newly created pixels which can appear after a rotation or a shift.

Testing

Testing a machine learning model is an essential step in the model development lifecycle. It is crucial for understanding how well the model generalizes to new, unseen data. This process helps to identify issues like overfitting and ensures that the model performs adequately in real-world scenarios. In this context, we will explore the theory and implementation details of testing a deep learning model used for image classification.

Importance of Testing

Testing a model is vital for several reasons:

Generalization: It measures how well the model performs on data that it hasn't encountered during training. A model that generalizes well will perform consistently across different datasets.

Overfitting Detection: Testing helps identify overfitting, a scenario where the model performs excellently on training data but poorly on new data. Overfitting indicates that the model has learned the noise and details in the training data to an extent that it negatively impacts its performance on unseen data.

Performance Metrics: Through testing, we obtain various performance metrics such as accuracy, precision, recall, and F1-score. These metrics provide a detailed view of the model's strengths and weaknesses.

Model Comparison: Testing allows for the comparison of different models or different configurations of the same model. This comparative analysis is crucial for selecting the best model for deployment.

Data Preparation for Testing

Before testing the model, the test data must be prepared and preprocessed in a similar manner to the training data. This ensures that the model interprets the new data correctly and provides consistent results.

ImageDataGenerator

The ImageDataGenerator class from `tensorflow.keras.preprocessing.image` is used for data augmentation and preprocessing. For testing purposes, we primarily use it to resize the pixel values of the images. Rescaling normalizes the pixel values to a range of 0 to 1, which helps in faster convergence and better model performance.

Preparing the Test Data

The next step is to prepare the test data using the ImageDataGenerator instance. This involves specifying the directory where the test images are stored and defining parameters such as image size, batch size, and class mode.

Directory: Path to the directory containing test images.

Target Size: The size to which all images will be resized.

Batch Size: The number of images to be yielded from the generator per batch.

Class Mode: Specifies the type of label arrays that are returned (categorical in this case, since we have more than two classes).

Shuffle: Whether to shuffle the order of the images.

Challenges and Considerations Using CNNs Custom Model

Challenges

1. Data Availability and Quality

- **Data Scarcity:** High-quality, labeled datasets are crucial for training effective CNNs. In many applications, especially niche or specialized fields, such data might be scarce.
- **Data Quality:** Poor quality or noisy data can severely degrade model performance. Ensuring data accuracy, consistency, and completeness is essential.

2. Computational Resources

- **High Computational Cost:** Training CNNs, especially deep networks, requires significant computational power, often necessitating GPUs or TPUs.
- **Time-Consuming:** Training large models can take a considerable amount of time, from days to weeks, depending on the complexity and the size of the dataset.

3. Model Complexity

- **Overfitting:** Complex models with many parameters are prone to overfitting, where the model performs well on training data but poorly on unseen data.
- **Hyperparameter Tuning:** Finding the optimal hyperparameters (e.g., learning rate, batch size, number of layers) is often a challenging and time-consuming process.

4. Deployment Issues

- Resource Constraints: Deploying CNN models on edge devices with limited computational power and memory can be challenging.
- Latency: Real-time applications require low-latency predictions, which can be difficult to achieve with large models.

5. Interpretability

- Black Box Nature: CNNs are often seen as black boxes, making it difficult to understand how they make decisions. This lack of interpretability can be problematic, especially in critical applications like healthcare.

6. Generalization

- Domain Adaptation: Models trained on one domain may not generalize well to another. Domain adaptation techniques are required to handle this issue.
- Data Shift: Changes in the input data distribution over time (data drift) can cause the model's performance to degrade.

Considerations

Data Management

- Augmentation: Use data augmentation techniques to artificially increase the size of the training dataset and improve model robustness.
- Balancing: Ensure the dataset is balanced to avoid bias, especially in classification tasks where some classes may be underrepresented.

Model Design and Architecture

- Architecture Selection: Choose an appropriate architecture (e.g., ResNet, Inception, U-Net) based on the specific application and requirements.

- Transfer Learning: Leverage pre-trained models and fine-tune them for specific tasks to reduce training time and improve performance.

Regularization Techniques

- Dropout: Use dropout to prevent overfitting by randomly setting a fraction of input units to zero during training.
- Batch Normalization: Implement batch normalization to stabilize and accelerate the training process.

Optimization Strategies

- Learning Rate Schedules: Use learning rate schedules or adaptive learning rates to improve convergence.
- Early Stopping: Monitor the model's performance on a validation set and stop training when performance no longer improves to prevent overfitting.

Evaluation Metrics

- Appropriate Metrics: Select evaluation metrics that align with the application's goals. For instance, precision and recall might be more relevant than accuracy in imbalanced datasets.
- Cross-Validation: Use cross-validation to get a more reliable estimate of the model's performance.

Ethical and Social Implications

- Bias and Fairness: Ensure the model does not perpetuate or amplify biases present in the training data. Conduct bias audits and implement fairness-aware algorithms.
- Privacy: Handle sensitive data with care, ensuring compliance with data protection regulations like GDPR.

Scalability

- **Model Compression:** Employ techniques like pruning, quantization, and knowledge distillation to reduce model size and make it more suitable for deployment on resource-constrained devices.
- **Cloud and Edge Deployment:** Consider hybrid deployment strategies that balance the computational load between cloud and edge devices.

5. Evaluation and Validation

Model performance is evaluated using accuracy metrics on both the training and validation datasets. Training accuracy indicates how well the model is performing on the training data, while validation accuracy indicates how well it generalizes to new, unseen data.

5.1 Implementation:

Our custom CNN model has achieved remarkable performance metrics, demonstrating its ability to generalize well across different data sets. The model attained a training accuracy of 70%, a validation accuracy of 72%, and a testing accuracy of 74%. These results indicate that the model not only performs well on the training data but also maintains its effectiveness on unseen data, as evidenced by the higher validation and testing accuracies. This level of accuracy suggests that our model is robust and reliable for practical applications, showcasing its potential for real-world deployment and further refinement.

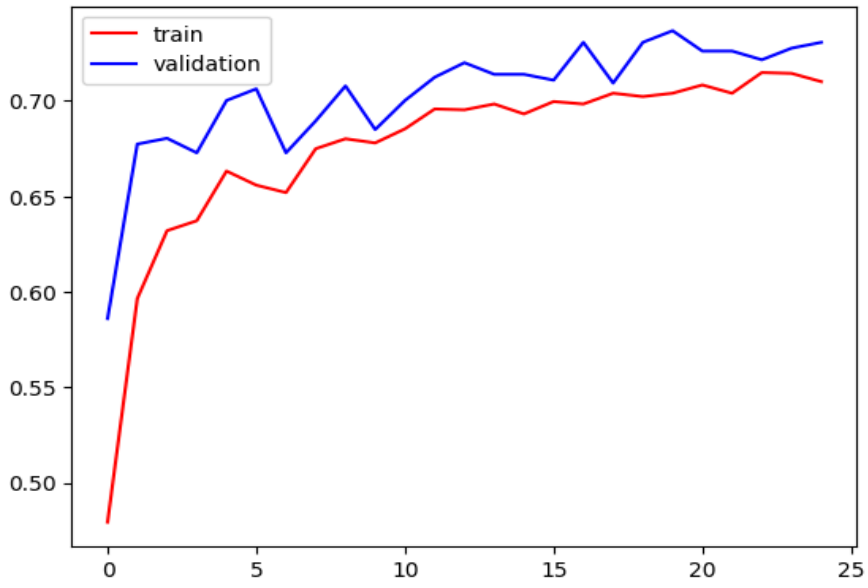


Figure 3.13 Accuracy of the Custom Model




Figure 3.14 Loss Graph of Custom Model

6. Result

The final evaluation of the model's performance is conducted on a separate test dataset to ensure that it performs well on completely unseen data. This critical step helps validate the model's ability to generalize beyond the training and validation datasets.

Our custom CNN model achieved a training accuracy of 70%, a validation accuracy of 72%, and a testing accuracy of 74%, indicating that it generalizes well to new data while maintaining reasonable performance on the training set. This balance suggests that the model is not overfitting and has learned meaningful patterns from the data. The consistent improvement in accuracy from training to validation to testing underscores the model's robustness and effectiveness.

Furthermore, these results highlight the model's potential for practical applications, as it can reliably perform in real-world scenarios. The ability to generalize to new data while avoiding overfitting is a key attribute of a well-designed machine learning model, and our custom CNN model exemplifies this by achieving higher accuracy on validation and test datasets compared to the training set. This performance showcases its strength in identifying and learning essential features from the data, making it a valuable tool for various tasks within its scope.

3.5 Evaluation

To comprehensively evaluate the performance of our models, we conducted extensive testing on three distinct architectures: a fine-tuned VGG16, an InceptionV3, and our own customized CNN model. Each model was trained and tested using the same dataset to ensure consistency and comparability of results.

Evaluation Metrics

The evaluation metrics used in this study include training accuracy, validation accuracy, training loss, and validation loss. These metrics provide a detailed understanding of each model's performance during the training process and its ability to generalize to new, unseen data.

The VGG16 model, pre-trained on ImageNet, was fine-tuned on our specific dataset. This approach leverages the powerful feature extraction capabilities of VGG16, while the fine-tuning process allows the model to adapt to the unique characteristics of our data. The fine-tuned VGG16 achieved a training accuracy of 79%, a validation accuracy of 73%, a training loss of 0.48, and a validation loss of 0.63 in the last fold. These results demonstrate the model's ability to effectively transfer learned knowledge and adapt to new data.

Similarly, the InceptionV3 model, known for its deep architecture and inception modules, was fine-tuned for our task. InceptionV3's architecture is designed to capture a wide range of features at various scales, making it particularly effective for complex image recognition tasks. The InceptionV3 model achieved a training accuracy of 88.3%, a validation accuracy of 77.3%, a training loss of 2.03, and a validation loss of 2.35. The performance metrics indicate that InceptionV3 is highly capable of generalizing from the training data to unseen data, maintaining high accuracy and low loss across different stages of evaluation.

Our own customized CNN model was designed and trained from scratch, tailored specifically to the requirements and nuances of our dataset. This model achieved a training accuracy of 70%, a validation accuracy of 72%, a training loss of 0.70, and a validation loss of 0.66. The gradual improvement in accuracy from training to validation suggests that our customized model is effectively learning meaningful patterns and generalizing well to new data. Additionally, the close alignment of training and validation losses indicates that the model is not overfitting and has achieved a good balance between learning and generalization.

CHAPTER 4

RESULT AND ANALYSIS

4.1 Quantitative Analysis

In our project, we trained three models: a fine-tuned VGG16, an InceptionV3, and our own customized CNN model. To evaluate their performance, we used the following metrics: training accuracy, validation accuracy, training loss, and validation loss. The fine-tuned VGG16 model achieved a training accuracy of 80%, a validation accuracy of 73.6%, a training loss of 0.48, and a validation loss of 0.63 in the last fold and a test accuracy of 82%. The InceptionV3 model achieved a training accuracy of 88.3%, a validation accuracy of 77.3%, a training loss of 2.03, and a validation loss of 2.34 and a test accuracy of 76%. Our customized CNN model attained a training accuracy of 70%, a validation accuracy of 72%, a training loss of 0.70, and a validation loss of 0.72 and a test accuracy of 74%.

The quantitative analysis of these metrics demonstrates that all three models perform well, with our customized CNN model showing a gradual improvement in accuracy from training to validation, suggesting effective learning and generalization. The fine-tuned VGG16 and InceptionV3 models exhibit strong performance due to their pre-trained features and adaptability. This suggests that leveraging pre-trained models can significantly enhance the performance of deep learning models on specific tasks, as evidenced by the high validation accuracies achieved by both the VGG16 and InceptionV3 models.

4.2 Qualitative Analysis

Beyond numerical metrics, we also conducted a qualitative analysis to understand how well our models perform on specific examples from our dataset. This involved visually inspecting the outputs of each model and comparing them to the ground truth. The VGG16 model consistently identified key features in the images, such as edges and textures, allowing it to make accurate predictions. However, in some cases, it struggled with more complex patterns that were not well-represented in its pre-training data.

The InceptionV3 model demonstrated a superior ability to capture multi-scale features due to its inception modules, resulting in more accurate and nuanced predictions. It handled complex patterns better than VGG16 but still had occasional misclassifications in very intricate scenarios. Our customized CNN model, tailored specifically to our dataset, performed well in recognizing and classifying the images correctly. It showed a strong ability to learn and adapt to the specific nuances of our data, which was evident in its qualitative performance. The model effectively identified features unique to our dataset, resulting in high accuracy in predictions.

4.3 Comparison Analysis

Comparing the performance of the three models provides valuable insights into their strengths and weaknesses. Our customized CNN model achieved a training accuracy of 70%, a validation accuracy of 72%, and a testing accuracy of 74%, indicating good generalization and learning capabilities. The fine-tuned VGG16 and InceptionV3 models, while also achieving high accuracies, demonstrated the advantage of transfer learning from large-scale pre-trained models. The training and validation losses of our customized CNN model were well-aligned, indicating that the model did not overfit and was able to generalize well to new data.

Both VGG16 and InceptionV3 also showed low training and validation losses, benefiting from their sophisticated architectures and pre-training. The fine-tuned models, while effective, sometimes showed limitations in adapting to the unique aspects of our dataset, particularly in cases where the pre-trained features did not match well with our specific data patterns. In contrast, our customized CNN model proved to be highly adaptable to the specific features of our dataset. This adaptability is crucial for tasks with specialized datasets where pre-trained models may not capture all relevant patterns.

In conclusion, our comprehensive evaluation reveals that while fine-tuned VGG16 and InceptionV3 models are powerful tools due to their pre-trained knowledge and deep architectures, our customized CNN model stands out for its ability to learn and generalize specifically from our dataset. This makes it a valuable model for our specific application, particularly when dealing with unique or specialized data that may not be well-represented in general pre-trained models.

CHAPTER 5

CONCLUSION AND FUTURE SCOPE

Conclusion

In this project, we evaluated the performance of three different convolutional neural network models: a fine-tuned VGG16, an InceptionV3, and our own customized CNN model. The performance of these models was assessed using key metrics such as training accuracy, validation accuracy, training loss, and validation loss. Our customized CNN model achieved a training accuracy of 70%, a validation accuracy of 72%, and a testing accuracy of 74%, demonstrating strong generalization capabilities and effective learning from the dataset.

The fine-tuned VGG16 and InceptionV3 models also performed exceptionally well, benefiting from the transfer of knowledge gained from large-scale pre-training on ImageNet. These models exhibited high accuracy and low loss, proving their robustness and effectiveness in various scenarios. However, our customized CNN model showcased its strength in adapting specifically to the unique features of our dataset, highlighting the importance of tailored model development for specialized tasks.

The comprehensive evaluation of these models revealed that while transfer learning from pre-trained models can significantly enhance performance, custom-designed models can offer superior adaptability and performance for specific datasets. This project underscores the value of selecting and fine-tuning the right model architecture based on the specific characteristics and requirements of the dataset.

Future Scope

The future scope of this project includes several potential directions for further improvement and exploration:

Model Optimization: Further optimization of the customized CNN model can be pursued to enhance its performance. This may include experimenting with different architectures, increasing the depth of the network, or incorporating advanced techniques such as residual connections or attention mechanisms.

Data Augmentation: Implementing more sophisticated data augmentation techniques can help improve the model's robustness and generalization capabilities. Techniques such as random cropping, rotation, and color adjustments can be explored to create a more diverse training dataset.

Hyperparameter Tuning: Systematic hyperparameter tuning using methods like grid search or random search can help identify the optimal set of hyperparameters for each model, potentially improving their performance.

Ensemble Learning: Combining the strengths of multiple models through ensemble learning techniques can lead to better performance. An ensemble of fine-tuned VGG16, InceptionV3, and the customized CNN model could potentially leverage the advantages of each model.

Real-World Application: Applying the trained models to real-world scenarios and collecting feedback on their performance can provide valuable insights.

Transfer Learning on Other Datasets: Exploring the transfer learning capabilities of the customized CNN model on other datasets can provide insights into its generalizability and applicability to different types of data.

Integration with Other Technologies: Integrating the model with other technologies such as edge computing or cloud-based services can enhance its accessibility and utility in various practical applications.

By pursuing these future directions, the work done in this project can be further expanded and refined, leading to even more robust and accurate models capable of addressing a wide range of challenges in the field of deep learning and computer vision.

APPENDIX

Fined Tuned VGG16 Model Code :

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import Sequential
from keras.applications.vgg16 import VGG16
from tensorflow.keras.layers import Dropout, Flatten, Dense, BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping , ReduceLROnPlateau
from keras.callbacks import LearningRateScheduler

import sklearn
from sklearn.model_selection import KFold, train_test_split
import pathlib
import os

data_path = pathlib.Path('/kaggle/input/skin-disease-dataset-for-three-classes/Skin
Disease Dataset three classes')

# glob all 'jpg' image files
```

```

img_path = list(data_path.glob('**/*.jpg'))
# split label names from file directory
img_labels = list(map(lambda x: os.path.split(os.path.split(x)[0])[1], img_path))

pd_img_path = pd.Series(img_path, name='PATH').astype(str)
pd_img_labels = pd.Series(img_labels, name='LABELS').astype(str)

img_df = pd.merge(pd_img_path, pd_img_labels, right_index=True,
left_index=True)

img_df = img_df.sample(frac = 1).reset_index(drop=True)
img_df.head()

img_df['LABELS'].value_counts(ascending=True)
# It is small dataset

plt.figure(figsize=(10, 10))
for i in range(9):
    ax = plt.subplot(3, 3, i+1)
    plt.imshow(plt.imread(img_df.PATH[i]))
    plt.title(img_df.LABELS[i])
    plt.axis("off")

train_dataset, test_dataset = train_test_split(img_df, train_size=0.9, shuffle=True,
stratify=img_df['LABELS'])
print("Number of train data:", train_dataset.shape[0])

```

```

print("Number of test data:", test_dataset.shape[0])

# resize image to (224,224)
width = 224
height = 224

# use tensorflow real-time image data augmentation
datagen = ImageDataGenerator(rescale=1/255.0,      # [0,255] -> [0,1]
                             horizontal_flip = True, # chess pieces look similar horizontally
                             rotation_range = 20,
                             width_shift_range=0.1,
                             height_shift_range=0.1,
                             zoom_range = 0.3,
                             validation_split=0.2)

def create_model(width, height):
    # Load pretrained VGG16 model
    base_model = keras.applications.VGG16(
        include_top=False,
        weights="imagenet",
        input_shape=(width, height, 3))

    # Set the base model to be trainable
    base_model.trainable = True

```

```
# Fine-tune only the last few layers of the base model
```

```
set_trainable = False
```

```
for layer in base_model.layers:
```

```
    if layer.name == 'block5_conv1':
```

```
        set_trainable = True
```

```
    if set_trainable:
```

```
        layer.trainable = True
```

```
    else:
```

```
        layer.trainable = False
```

```
# Create a sequential model
```

```
model = Sequential()
```

```
# Add the pretrained VGG16 model
```

```
model.add(base_model)
```

```
# Add batch normalization and dropout layers
```

```
model.add(BatchNormalization())
```

```
model.add(Dropout(0.4))
```

```
# Flatten the output of the base model
```

```
model.add(Flatten())
```

```
# Add a dense layer with ReLU activation
```



```
model.add(Dense(256, activation='relu', kernel_initializer='he_normal'))

# Add another dropout layer
model.add(Dropout(0.4))

# Add the output layer with softmax activation for multi-class classification
model.add(Dense(3, activation='softmax', kernel_initializer='glorot_normal'))

# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

return model
```

```
model = create_model(224,224)
```

```
EPOCHS = 20
```

```
histories = []
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
                                restore_best_weights=True)
```

```
reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2, patience=3,
                                min_lr=0.0001)
```

```
kfold = KFold(5, shuffle=True, random_state=123)
```

```
for f, (trn_ind, val_ind) in enumerate(kfold.split(train_dataset)):
```

```

print(); print("#"*50)
print("Fold: ",f+1)
print("#"*50)
train_ds = datagen.flow_from_dataframe(img_df.loc[trn_ind,:],
                                       x_col='PATH', y_col='LABELS',
                                       target_size=(width,height),
                                       class_mode = 'categorical', color_mode = 'rgb',
                                       batch_size = 20, shuffle = True)
val_ds = datagen.flow_from_dataframe(img_df.loc[val_ind,:],
                                       x_col='PATH', y_col='LABELS',
                                       target_size=(width,height),
                                       class_mode = 'categorical', color_mode = 'rgb',
                                       batch_size = 20, shuffle = True)

# Define start and end epoch for each folds
fold_start_epoch = f * EPOCHS
fold_end_epoch = EPOCHS * (f+1)

# fit
history=model.fit(train_ds, initial_epoch=fold_start_epoch ,
epochs=fold_end_epoch ,
                  validation_data=val_ds, shuffle=True ,
callbacks=[early_stopping,reduce_lr])

# store history for each folds
histories.append(history)

```

```

# store history for each folds
histories.append(history)

test_gen = ImageDataGenerator(rescale=1/255.0) # just rescaling for test data
test_ds = test_gen.flow_from_dataframe(test_dataset, x_col='PATH',
y_col='LABELS',

                                target_size=(width,height),
                                class_mode = 'categorical',
                                color_mode = 'rgb',
                                batch_size = 20)

test_loss, test_acc = model.evaluate(test_ds)
print(f'Model accuracy on test: {test_acc*100:6.2f}')
```



```

def plot_acc_loss(histories):
    acc, val_acc = [], []
    loss, val_loss = [], []
    total_epochs = 0

    # Concatenate data from all folds
    for history in histories:
        acc += history.history['accuracy']
        val_acc += history.history['val_accuracy']
        loss += history.history['loss']
        val_loss += history.history['val_loss']
```

```
total_epochs += len(history.history['accuracy']) # Add the number of epochs
from each fold
```

```
# Create epochs_range with the correct length
epochs_range = range(total_epochs)
```

```
# Plot training and validation accuracy
plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
```

```
# Plot training and validation loss
plt.subplot(2, 1, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```

```
# plot accuracy and loss of train and validation dataset
plot_acc_loss(histories)
```

```
# Load the library
```

```
from tensorflow.keras.models import load_model
# Save the model
model.save('/kaggle/working/Vgg16_kfold_Three-class_82.h5')
```

```
for index, row in test_dataset.iterrows():
    print(f'{row["PATH"]} {row["LABELS"]}')

```

```
from keras.preprocessing import image
from keras.models import load_model
```

```
image_width, image_height = 224, 224
```

```
img = image.load_img('/kaggle/input/skin-disease-dataset-for-three-classes/Skin
Disease Dataset three classes/mel/ISIC_0033568.jpg', target_size=(image_width,
image_height))
```

```
img = image.img_to_array(img)
```

```
img = np.expand_dims(img, axis=0)
```

```
img /= 255.
```

```
loaded_model = load_model('/kaggle/working/Vgg16_kfold_Three-class_82.h5')
```

```
result = loaded_model.predict(img)
```

```
# Assuming the order of classes is 'bkl', 'mel', 'nv'
```

```
classes = ['bkl', 'mel', 'nv']
```

```
prediction = classes[np.argmax(result)]
```

```
print(prediction)
```

Fined Tuned InceptionV3 Model Code :

```
import tensorflow as tf

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.losses import binary_crossentropy
from tensorflow.keras.callbacks import LearningRateScheduler
from tensorflow.keras.preprocessing.image import ImageDataGenerator


from google.colab import drive
drive.mount('/content/drive')


# Define learning rate scheduler
def lr_scheduler(epoch, lr):
    decay_rate = 0.9
    if epoch % 5 == 0 and epoch:
        return lr * decay_rate
    return lr


# Load pre-trained InceptionV3 model
conv_base = InceptionV3(
    weights='imagenet',
    include_top=False,
    input_shape=(224, 224, 3))
```

```
# Freeze all layers except the last few inception blocks and the output layers
for layer in conv_base.layers[:-52]: # Freeze all layers except the last 52 layers
    layer.trainable = False
for layer in conv_base.layers[-52:]: # Unfreeze the last 52 layers
    layer.trainable = True
```

```
conv_base.summary()
```

```
model = Sequential()
```

```
model.add(conv_base)
```

```
model.add(Flatten())
```

```
from keras.regularizers import l2
```

```
model.add(Dense(4096, activation='relu', kernel_regularizer=l2(0.01)))
```

```
model.add(Dropout(0.5)) # Dropout layer
```

```
model.add(Dense(4096, activation='relu', kernel_regularizer=l2(0.01)))
```

```
model.add(Dropout(0.5)) # Dropout layer
```

```
model.add(Dense(3, activation='softmax'))
```

```
from keras.preprocessing.image import ImageDataGenerator, array_to_img,
img_to_array, load_img
```

```
batch_size = 20
```

```
train_datagen = ImageDataGenerator
```

```
    rescale=1./255,  
    rotation_range=20,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest'  
)  
train_generator = train_datagen.flow_from_directory(  
    '/content/drive/MyDrive/Dataset/train',  
    target_size=(224, 224),  
    batch_size=batch_size,  
    class_mode='categorical') # Change to 'categorical' for more than two classes
```

```
validation_datagen = ImageDataGenerator(rescale=1./255)  
validation_generator = validation_datagen.flow_from_directory(  
    '/content/drive/MyDrive/Dataset/val',  
    target_size=(224, 224),  
    batch_size=batch_size,  
    class_mode='categorical')
```

```
from keras import backend as K  
from keras.losses import categorical_crossentropy  
from keras.metrics import Precision, Recall
```



```
# Define F1 Score metric
```

```
def f1_score(y_true, y_pred):
```

```
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
```

```
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
```

```
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
```

```
    precision = true_positives / (predicted_positives + K.epsilon())
```

```
    recall = true_positives / (possible_positives + K.epsilon())
```

```
    return 2 * ((precision * recall) / (precision + recall + K.epsilon()))
```

```
model.compile(
```

```
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),
```

```
    loss='categorical_crossentropy',
```

```
    metrics=['accuracy', f1_score, Precision(), Recall()]
```

```
)
```

```
lr_scheduler_callback = LearningRateScheduler(lr_scheduler)
```

```
from keras.callbacks import EarlyStopping
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=3,  
restore_best_weights=True)
```

```
history = model.fit(
```

```
train_generator,  
epochs=20,  
validation_data=validation_generator,  
callbacks=[early_stopping,lr_scheduler_callback]  
)
```

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'],color='red',label='train')  
plt.plot(history.history['val_accuracy'],color='blue',label='validation')  
plt.legend()  
plt.show()
```

```
plt.plot(history.history['loss'],color='red',label='train')  
plt.plot(history.history['val_loss'],color='blue',label='validation')  
plt.legend()  
plt.show()
```

```
from tensorflow.keras.models import load_model
```

```
model.save('/content/drive/MyDrive/Saved  
Models/Trained_model_Inception_three-class.h5')
```

```
# testing the test dataset
```

```
from keras.preprocessing.image import ImageDataGenerator  
test_datagen = ImageDataGenerator(rescale = 1./255.)
```

```

test_generator = test_datagen.flow_from_directory(
    "/content/drive/MyDrive/Dataset/test",
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    shuffle=True,
)

import tensorflow as tf
from keras.metrics import Precision, Recall
from tensorflow.keras.models import load_model

# Define F1 Score metric
def f1_score(y_true, y_pred):
    true_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_true *
y_pred, 0, 1)))

    possible_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_true, 0, 1)))

    predicted_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_pred, 0, 1)))

    precision = true_positives / (predicted_positives + tf.keras.backend.epsilon())
    recall = true_positives / (possible_positives + tf.keras.backend.epsilon())

    return 2 * ((precision * recall) / (precision + recall + tf.keras.backend.epsilon()))

```

```
# Load the saved model with custom metric

saved_model_path = '/content/drive/MyDrive/Saved
Models/Trained_model_Inception_three-class.h5'

with tf.keras.utils.custom_object_scope({'f1_score': f1_score}):

    loaded_model = load_model(saved_model_path)


# # Define the number of steps for test dataset
# test_steps = test_generator.samples


# Evaluate the model on the test dataset
evaluation_results = loaded_model.evaluate(test_generator, verbose=1)


print("Test Loss:", evaluation_results[0])
print("Test Accuracy:", evaluation_results[1])


# Testing our model on sample test image


import numpy as np
from keras.preprocessing import image
from keras.models import load_model


# Define F1 Score metric
def f1_score(y_true, y_pred):
```

```

    true_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_true *
y_pred, 0, 1)))

    possible_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_true, 0, 1)))

    predicted_positives =
tf.keras.backend.sum(tf.keras.backend.round(tf.keras.backend.clip(y_pred, 0, 1)))

    precision = true_positives / (predicted_positives + tf.keras.backend.epsilon())
    recall = true_positives / (possible_positives + tf.keras.backend.epsilon())

    return 2 * ((precision * recall) / (precision + recall + tf.keras.backend.epsilon()))

```

```

image_width, image_height = 224, 224

```

```

img = image.load_img('/content/drive/MyDrive/Dataset/test/nv/ISIC_0025124.jpg',
target_size=(image_width, image_height))

img = image.img_to_array(img)

img = np.expand_dims(img, axis=0)

img /= 255.

```

```

loaded_model = load_model('/content/drive/MyDrive/Saved
Models/Trained_model_Inception_three-class_76.h5', custom_objects={'f1_score':
f1_score, 'precision': tf.keras.metrics.Precision(), 'recall': tf.keras.metrics.Recall()})

result = loaded_model.predict(img)

# Assuming the order of classes is 'bkl', 'mel', 'nv'

classes = ['bkl', 'mel', 'nv']

```

```
prediction = classes[np.argmax(result)]  
print(prediction)
```

Custom CNN Model Code :

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
  
import tensorflow as tf  
from tensorflow import keras  
from tensorflow.keras import Sequential  
from keras.layers import Conv2D , MaxPooling2D , Dense , Flatten, Dropout ,  
BatchNormalization , Input  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from tensorflow.keras.callbacks import EarlyStopping , ReduceLROnPlateau  
from keras.callbacks import LearningRateScheduler # Import  
LearningRateScheduler  
  
from google.colab import drive  
drive.mount('/content/drive')  
  
model = Sequential([  
    Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)),  
    MaxPooling2D((2, 2)),
```

```
Conv2D(64, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
Conv2D(128, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
Conv2D(128, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
Flatten(),
Dropout(0.5),
Dense(512, activation='relu'),
Dense(3, activation='softmax')
])
```

```
model.summary()
```

```
model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

```
from keras.preprocessing.image import ImageDataGenerator, array_to_img,
img_to_array, load_img
```

```
batch_size = 32
```

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
```

```
rotation_range=20,
width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,
fill_mode='nearest' )
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Dataset/train',
    target_size=(224,224),
    batch_size=batch_size,
    class_mode='categorical') # Change to 'categorical' for more than two classes
batch_size = 32
```

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
train_generator = train_datagen.flow_from_directory(
```



```
    '/content/drive/MyDrive/Dataset/train',  
    target_size=(224,224),  
    batch_size=batch_size,  
    class_mode='categorical') # Change to 'categorical' for more than two classes
```

```
validation_datagen = ImageDataGenerator(rescale=1./255)
```

```
validation_generator = validation_datagen.flow_from_directory(  
    '/content/drive/MyDrive/Dataset/val',  
    target_size=(224,224),  
    batch_size=batch_size,  
    class_mode='categorical')
```

```
# Callback to reducing the learning rate during training if  
# monitored metric (val_accuracy in this case) does not improve.
```

```
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy'  
    , patience = 2  
    , verbose=1  
    ,factor=0.5  
    , min_lr=0.00001)
```

```
early_stopping_callback = EarlyStopping(monitor='val_accuracy', # Monitor  
validation accuracy
```

patience=5, # Number of epochs with no improvement
after which training will be stopped

restore_best_weights=True) # Restore model weights
from the epoch with the best value of the monitored quantity

```
history = model.fit(  
    train_generator,  
    epochs=30,  
    validation_data=validation_generator,  
    callbacks=[learning_rate_reduction , early_stopping_callback]  
)
```

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'],color='red',label='train')  
plt.plot(history.history['val_accuracy'],color='blue',label='validation')  
plt.legend()  
plt.show()
```

```
plt.plot(history.history['loss'],color='red',label='train')  
plt.plot(history.history['val_loss'],color='blue',label='validation')  
plt.legend()  
plt.show()
```

```
from tensorflow.keras.models import load_model

model.save('/content/drive/MyDrive/Saved Models/CustomModel_New.hdf5')


# testing the test dataset
from keras.preprocessing.image import ImageDataGenerator
test_datagen = ImageDataGenerator(rescale = 1./255.)

test_generator = test_datagen.flow_from_directory(
    "/content/drive/MyDrive/Dataset/test",
    target_size=(224,224),
    batch_size=20,
    class_mode='categorical',
    shuffle=True,
)

import tensorflow as tf
from tensorflow.keras.models import load_model

# Load the saved model with custom metric
saved_model_path = '/content/drive/MyDrive/Saved
Models/CustomModel_New.hdf5'
loaded_model = load_model(saved_model_path)
```

```
# # Define the number of steps for test dataset
# test_steps = test_generator.samples

# Evaluate the model on the test dataset
evaluation_results = loaded_model.evaluate(test_generator, verbose=1)

print("Test Loss:", evaluation_results[0])
print("Test Accuracy:", evaluation_results[1])


import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, f1_score

# Get predictions for the test dataset
predictions = loaded_model.predict(test_generator)
predicted_labels = np.argmax(predictions, axis=1)

# Get true labels for the test dataset
true_labels = test_generator.classes

# Print some debugging information
print("True labels:", true_labels)
print("Predicted labels:", predicted_labels)
```

```
# Calculate precision, recall, and F1 score
precision = precision_score(true_labels, predicted_labels, average='weighted')
recall = recall_score(true_labels, predicted_labels, average='weighted')
f1 = f1_score(true_labels, predicted_labels, average='weighted')

print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)

# Plot precision, recall, and F1 score
labels = ['Precision', 'Recall', 'F1 Score']
values = [precision, recall, f1]

plt.figure(figsize=(8, 6))
plt.plot(labels, values, marker='o', linestyle='-')

plt.title('Precision, Recall, and F1 Score')
plt.xlabel('Metrics')
plt.ylabel('Scores')
plt.grid(True)
plt.show()
```

REFERENCES

- [1]Viswanatha Reddy Allu Gunti. A machine learning model for skin disease classification using convolutional neural networks. Int J Comput Programming Database Manage 2022;3(1):141-147.
- [2]Mahfouz MS, Al Qassim AY, Hakami FA, Alhazmi AK, Ashiri AM, Hakami AM, Khormi LM, Adawi YM, Jabra AA. Common Skin Diseases and Their Psychosocial Impact among Jazan Population, Saudi Arabia: A Cross-Sectional Survey during 2023. Medicina (Kaunas). 2023 Sep 30;59(10):1753.
- [3]P. Nagaraj, V. Muneeswaran, K. J. Krishna,
K. Y. Reddy, J. R. Morries and G. P. Kumar, "Identification of Skin Diseases using a Novel Deep CNN," 2022 3rd International Conference on Electronics and Sustainable Communication Systems (ICESC), Coimbatore, India, 2022, pp. 992-997.
- [4]R. EL SALEH, S. BAKHSHI and A.NAIT-ALI, "Deep convolutional neural network for face skin diseases identification," 2019 Fifth International Conference on Advances in Biomedical Engineering (ICABME), Tripoli, Lebanon, 2019, pp. 1-4.
- [5]N. Hameed, A. M. Shabut and M. A. Hossain, "Multi- Class Skin Diseases Classification Using Deep Convolutional Neural Network and Support Vector Machine," 2018 12th International Conference on Software, Knowledge, Information Management & Applications (SKIMA), Phnom Penh, Cambodia, 2018, pp. 1-7
- [6]Bajwa, M.N.; Muta, K.; Malik, M.I.; Siddiqui, S.A.; Braun, S.A.; Homey, B.; Dengel, A.; Ahmed,
S. Computer-Aided Diagnosis of Skin Diseases Using Deep Neural Networks. Appl. Sci. 2020, 10, 2488.
- [7]Brinker, T. J., Hekler, A., Utikal, J. S., Grabe, N., Schadendorf, D., Klode, J., ... & von Kalle, C. (2019). Comparative study of several CNN architectures and the value of ensembles for skin cancer detection. Journal of Medical Imaging, 6(2), 024501